

Automatic prompt debiasing via contrastive reflection for fairness enhancement in large language models

Juncheng Hu ^{a,b}, Yingji Li ^{a,b,*}, Rui Song ^a, Kedi Lyu ^a, Yi Chang ^c

^a College of Computer Science and Technology, Jilin University, Changchun, 130012, China

^b Engineering Research Center for Network Technology and Applied Software, Ministry of Education, Jilin University, Changchun, 130012, China

^c School of Artificial Intelligence, Jilin University, Changchun, 130012, China

ARTICLE INFO

Keywords:

Large language models
Fairness
Social bias
Prompt optimization
Reflection mechanism

ABSTRACT

With the rapid development of Large Language Models (LLMs), increasing attention has been paid to the potential social biases in their outputs. Most existing approaches rely on manual prompt engineering to mitigate social biases in LLMs. However, manually designed task prompts are resource-intensive and heavily dependent on expert knowledge. More importantly, fixed prompt templates are difficult to adapt to diverse debiasing scenarios, limiting both flexibility and scalability. To this end, this paper proposes a self-debiasing framework based on automatic prompt engineering, Automatic Prompt Debiasing via contrastive reflection (called **APD**), which enables LLMs to improve their fairness autonomously. Specifically, we use an LLM as a prompt optimizer, leveraging its powerful inference and generation capabilities to automatically create prompts that reduce stereotypical outputs. Furthermore, we design a contrastive reflection mechanism to align LLMs with fairness objectives, enabling them to refine prompts by contrasting fair and biased features in example groups, thereby enhancing prompt optimization. Experiments on 14 LLMs, covering 9 singular biases and 2 intersectional biases across two benchmark datasets, demonstrate the superior debiasing effectiveness of APD in various scenarios and its potential to reduce the original bias of multiple LLMs to negligible levels. Additionally, generalization experiments across multiple dimensions validate the excellent scalability of APD.

1. Introduction

Large Language Models (LLMs) [1,2] have demonstrated unprecedented capabilities and, as a result, have been widely adopted in various real-world applications [3,4]. Numerous studies have shown that LLMs trained on large-scale web data tend to internalize social biases present in the data and generate human-like stereotypical responses [5–7]. Deploying LLMs with social biases in real-world settings may hinder social equity and result in unforeseen negative consequences [8–10]. For instance, LLMs exhibiting gender or racial biases may produce discriminatory outputs in high-stakes domains such as hiring assistance, legal question answering, or medical consultation, e.g., favoring male candidates, reinforcing harmful stereotypes, or overlooking symptoms prevalent among minority groups, thereby amplifying societal inequalities [11]. Effectively mitigating social biases in LLMs has thus become a critical and urgent problem in the field of data mining, particularly in ensuring that model-driven insights and decisions are fair, accountable, and socially responsible in practical applications.

* Corresponding author at: College of Computer Science and Technology, Jilin University, Changchun, 130012, China.

E-mail addresses: jchu@jlu.edu.cn (J. Hu), yingjili@jlu.edu.cn (Y. Li), songrui@jlu.edu.cn (R. Song), kdlyu@jlu.edu.cn (K. Lyu), yichang@jlu.edu.cn (Y. Chang).

<https://doi.org/10.1016/j.artint.2026.104625>

Received 14 August 2025; Received in revised form 14 September 2026; Accepted 14 September 2026

Available online 16 September 2026

0004-3702/© 2026 Elsevier B.V. All rights are reserved, including those for text and data mining, AI training, and similar technologies.

Due to their large parameter size and semi-open source nature, fine-tuning LLMs or accessing the probability distribution of model outputs is generally considered an inefficient strategy for mitigating social biases. Several recent works have attempted to employ manual prompt engineering [12,13] to debias LLMs. The objective is to design prompts that reduce biased responses. Oba et al. [14] mitigates potential stereotypes in LLMs by constructing textual preambles that include counterfactual statements derived from human-designed templates and real-world statistics. Some studies have shown that enhancing the reasoning capabilities of LLMs through Chain of Thought (CoT) [15,16] or Dual Process Theory [17,18] can also indirectly improve the fairness of their predictions. Kamruzzaman and Kim [17] also enable LLMs to mimic human cognitive processes by assigning specific roles to the model, thereby reducing stereotypes in their outputs.

Although these methods have achieved some success, they still exhibit notable limitations. On the one hand, hand-crafted task prompts are constrained by labor costs and external resources, and they heavily rely on the expertise of prompt engineers. On the other hand, fixed prompt templates struggle to provide personalized debiasing prompts across diverse tasks, resulting in limited flexibility and scalability. In addition, the debiasing performance of prompts based on CoT or person-assigned roles is often inconsistent and may even exacerbate biases in certain tasks [19,20].

To this end, we propose Automatic Prompt Debiasing via contrastive reflection (APD), a self-debiasing framework designed to mitigate social biases in LLM outputs and enhance response fairness. Unlike debiasing methods based on manual prompt engineering, APD leverages the reasoning and generation capabilities of LLMs to perform automatic self-debiasing, eliminating the need for manual customization and offering greater flexibility and scalability. Specifically, we employ an LLM-based automatic prompt optimization strategy to generate task-specific prompts that can mitigate the biased responses produced by LLMs. These prompts are iteratively refined to progressively improve their debiasing effectiveness. To better align LLMs with debiasing objectives, we propose a novel contrastive reflection mechanism. We guide LLMs to observe *contrast example groups* that contain both fair and biased features. This encourages deep reflection on the shortcomings of current prompts and facilitates the generation of insightful modification suggestions. The LLM-based prompt generator expands the current prompts under the combined guidance of contrast example groups and modification suggestions. A search strategy is then used to filter and select the best candidate as the final prompt for the current iteration.

Our main contributions are summarized as follows:

- We propose a novel self-debiasing framework based on automatic prompt engineering. It leverages the reasoning and self-reflection capabilities of LLMs to optimize task prompts and improve response fairness, eliminating the need for manual design and offering greater flexibility and scalability.
- We introduce a contrastive reflection mechanism to enhance the LLM-based automatic prompt optimization process. It encourages LLMs to identify both fair and biased features and align more effectively with the fairness objective through deeper reflection.
- Experiments on 14 LLMs involving 9 singular and 2 intersectional biases across two benchmark datasets verify that APD can reduce the original bias to 0% in some cases and significantly outperform manual prompt engineering methods. Furthermore, additional evaluations confirm that APD exhibits strong generalization capabilities.

2. Related work

2.1. Debiasing methods for PLMs

Debiasing strategies for Pre-trained Language Models (PLMs) are conventionally classified into two overarching paradigms: model-centric and data-centric approaches, distinguished by their respective optimization objectives and intervention mechanisms. Model-centric methods seek to mitigate bias by directly modifying the model's internal learning dynamics or representational structures. This line of work encompasses a range of techniques, including contrastive learning [6,21], projection-based debiasing in the representation space [22,23], adversarial training frameworks [24,25], and the incorporation of fairness-aware objective functions [26,27]. In parallel, data-centric approaches operate on the assumption that social biases originate primarily from the training data. Accordingly, these methods aim to mitigate bias by improving data quality or balance, utilizing techniques such as counterfactual data augmentation [7,28,29], instance reweighting [30,31], and data correction [32,33]. Beyond direct data modification, Schick et al. [34] showed that prompting can enable PLMs to self-diagnose and suppress undesirable behaviors during generation, providing an alternative means of leveraging model feedback for bias mitigation.

Despite their demonstrated efficacy in standard PLMs, extending these traditional debiasing paradigms to LLMs remains non-trivial. A primary challenge stems from the architectural and operational characteristics of LLMs, namely that most traditional debiasing methods presuppose access to internal parameters or the ability to manipulate output probability distributions, capabilities that are typically restricted due to the proprietary nature and immense scale of contemporary LLMs. Although studies such as Schick et al. [34] have explored prompt-based approaches for guiding models toward less harmful behavior, systematic and scalable debiasing methods for LLMs remain underexplored. Consequently, these constraints necessitate the development of novel model-agnostic debiasing frameworks that can operate under limited access assumptions while still ensuring fairness and accountability.

2.2. Debiasing methods based on prompt engineering

Many studies have found that the output of LLMs is seriously affected by the input instructions [35–37]. Thus, prompt engineering [38,39] has emerged as a more efficient approach to debiasing LLMs. These methods mitigate biases in LLMs' responses by designing




Fig. 1. The overall framework of APD illustrated by the gender bias of BBQ dataset. The initial task prompt is adopted for the first iteration. In each iteration, the fairness evaluation module firstly evaluates the current prompt, and sends the evaluation results of the training data into the contrastive reflection module for sampling to obtain contrast example groups. Then, the contrastive reflection module performs self-reflection and gives improvement suggestions. Finally, the prompt generation module optimizes the current prompt to obtain new candidate prompts, which are searched and selected by the prompt search module and the updated prompts will go to the next iteration.

prompts that suppress stereotypes in the models' output. Existing debiasing methods based on prompt engineering mainly use hand-crafted prompt templates. For example, Kamruzzaman and Kim [17] adopt the persona-assignment strategy [40] and prepend the template "Adopt the identity of [persona]." to the original task prompts. They find that LLMs reduce stereotypical responses when assigned a persona, as the models imitate human cognitive processes from the perspective of the assigned identity. They also utilize a prompt strategy inspired by Dual Process Theory [41,42], in which System 2 refers to slow, effortful, and analytical thinking, leading to more reliable decisions akin to human cognition. The corresponding prompt is "You answer questions slowly and thoughtfully. Your answers are effortful and reliable.". Oba et al. [14] propose a method to suppress gender bias by adding counterfactual textual preambles to prompts, which are generated from hand-crafted templates and real-world statistics. In addition, several studies have found that optimization methods designed for general tasks can also help alleviate social biases in LLMs. For instance, Chain of Thought (CoT) [43,44], which enhances LLMs' reasoning abilities, can reduce biases by encouraging step-by-step thinking using prompts like "Let's think step by step." [15,16].

However, current debiasing methods based on prompt engineering suffer from several limitations. First, hand-crafted templates are resource-intensive, and selecting the most effective template is often challenging for practitioners. Second, the prompts used in these methods are often fixed and simplistic, making it difficult to ensure their effectiveness and adaptability when applied to new tasks. This highlights the need for more flexible and extensible prompting strategies. Finally, their debiasing effectiveness may be limited. For example, while CoT shows some improvement, other studies have reported that persona-assignment strategies can increase toxicity and implicit bias in models such as ChatGPT [19,20].

In this study, we propose a debiasing method based on automatic prompt engineering, which leverages LLM-driven prompt optimization to generate task-specific prompts. Unlike previous approaches, our method requires neither manual customization nor prior information, significantly reducing labor costs and reliance on external resources. Moreover, APD is an adaptive process capable of automatically generating tailored prompts for diverse tasks, offering greater efficiency, flexibility, and scalability.

3. Task definition

For a given LLM $\mathcal{M}$, we consider a dataset $\mathcal{D} = \{(x_i, y_i)\}_{i=1}^n$ associated with a social bias-related task, which is split into a training set D_{train} and a test set D_{test} . Each sample x_i contains information about a demographic group (e.g., *male/female* for gender, *white/black* for race). Given a fairness evaluation function $s(\cdot)$, which calculates the fairness score for $\mathcal{M}$ based on the model's response $\hat{y}_i = \mathcal{M}(x_i, p)$ and the true label y_i , our objective is to identify the optimal prompt p^* in the discrete natural language space via prompt engineering. The prompt should maximize fairness on $\mathcal{D}$, trained on D_{train} and evaluated on D_{test} , ensuring that $\mathcal{M}$'s responses mitigate social bias. The formal definition is as follows:

$$p^* = \underset{p}{\operatorname{argmax}} \mathbb{E}_{(x_i, y_i) \sim D_{train}} [s(\mathcal{M}(x_i, p), y_i)]. \quad (1)$$

Previous debiasing methods typically rely on manually defined prompts. In contrast, APD employs LLM-based automatic prompt optimization to generate and search for the best prompt, guided by self-feedback from the LLM itself.

4. Methodology

Automatic prompt engineering can significantly reduce the labor-intensive efforts required in manual prompt engineering [45]. As a result, there has been growing interest in LLM-based automatic prompt optimization. These approaches employ LLMs as prompt optimizers to generate ideal task prompts for target models [46,47]. Research has shown that LLMs, when used as prompt optimizers, can achieve performance comparable to that of expert human prompt engineers, positioning them as a promising new research paradigm [48–50].

Inspired by these related works, we adopt LLM-based automatic prompt optimization to mitigate social biases in LLMs output responses. By leveraging the powerful reasoning and generative capabilities of LLMs, we iteratively update task prompts to enhance the effectiveness of debiasing. However, even though LLMs perform well in many tasks, it is challenging to enable LLMs to understand subtle concepts of fairness and bias in human society and to engage in deep reflection specific to debiasing. Therefore, we propose a contrastive reflection mechanism to improve the efficiency and effectiveness of alignment in debiasing tasks. It leverages the self-reflection abilities of LLMs to identify fairness and bias features in the data. By allowing the model to observe contrast example groups, the mechanism guides the LLMs to generate constructive reflection suggestions tailored to specific debiasing tasks, thus strengthening the prompt optimizer's ability to optimize tasks.

Our proposed framework for automatic prompt debiasing via contrastive reflection, APD, is illustrated in Fig. 1. It consists of four main modules: the Fairness Evaluation Module, the Contrastive Reflection Module, the Prompt Generation Module, and the Prompt Search Module. During each round of iterative training, the fairness evaluation module first evaluates the current task prompt and sends the evaluation results to the contrastive reflection module. This module samples contrast example groups based on the evaluation, performs self-reflection, and provides improvement suggestions based on both the current prompt and the contrast example groups. The prompt generation module then consolidates all available information to optimize the prompt, generating new candidate prompts, which are subsequently selected and updated by the prompt search module to determine the final prompt for that iteration.

The overall algorithm is given in Algorithm 1. For a given task and bias category, APD takes an initial prompt set, a fairness evaluation function, and training data as inputs, and automatically optimizes candidate prompts through contrastive reflection, prompt generation, and prompt search. The process is fully automated and can be applied separately to different tasks or bias categories, with the optimized prompts reused for inference.

Algorithm 1: Overall algorithm of APD.

Require: LLM $\mathcal{M}$, training set D_{train} , fairness evaluation function $s(\cdot)$, initial set of prompts $\mathcal{P}_0$, iteration number T , mini-batch size m , number of a contrast example group K , suggestions number L

- 1: $\mathcal{P}^0 \leftarrow \{p_0\}$
- 2: **for** $t \leftarrow 0$ to $T - 1$ **do**
- 3: **for all** $p \in \mathcal{P}^t$ **do**
- 4: Sample mini-batch $D^m \subset D_{train}$
- 5: Feed D^m into $\mathcal{M}$ and obtain $\mathcal{Y}^m = \{\hat{y}_i\}_{i=1}^m$
- 6: Utilize $s(\cdot)$ to evaluate p and obtain $D^f = \{(x_i^f, y_i^f, \hat{y}_i^f)\}$ and $D^b = \{(x_i^b, y_i^b, \hat{y}_i^b)\}$
- 7: Sample contrast example group $\mathcal{E} = \{(e_i^f, e_i^b)\}_{i=1}^K \subset D^f \cup D^b$
- 8: Obtain suggestions $\mathcal{R} = \{r_i\}_{i=0}^L$ reflected by $\mathcal{M}$ based on p and $\mathcal{E}$
- 9: Generate new candidate prompt $\hat{p}$ based on p , $\mathcal{E}$, and $\mathcal{R}$
- 10: **end for**
- 11: $\hat{\mathcal{P}}^t \leftarrow \mathcal{P}^t \cup \hat{\mathcal{P}}^t$
- 12: $\hat{\mathcal{P}}_0^t \leftarrow \hat{\mathcal{P}}^t$
- 13: **for** $\theta \leftarrow 0$ to $\Theta - 1 = \log 2^{|\hat{\mathcal{P}}^t|} - 1$ **do**
- 14: Sample $D^{n_k} \subset D_{train}$
- 15: Evaluate $\forall \hat{p} \in \hat{\mathcal{P}}_0^t$ with $s(\hat{p}, D^{n_k})$
- 16: Filter out prompts with a bias score worse than the average score
- 17: $\hat{\mathcal{P}}_{\theta+1}^t \leftarrow \hat{\mathcal{P}}_\theta^t$
- 18: **end for**
- 19: $\mathcal{P}^{t+1} \leftarrow \hat{\mathcal{P}}_{\Theta-1}^t$
- 20: **end for**
- 21: **return** $\mathcal{P}^{T-1}$

4.1. Prompt initialization

For a given number of T iterations of training, when $t = 0 \in [0, T)$, an initial set of task prompts is required to begin the training. Existing research explores two main approaches for prompt initialization: *manual initialization*, which involves human-written task instructions [51,52], and *inductive initialization*, where LLMs generate initial task instructions based on few-shot examples [53,54]. To facilitate efficient and accurate optimization, we use manual optimization to provide an initial prompt, denoted as p_0 . It is important

to note that this approach differs from manual prompt engineering, as we only need to write basic task instructions without any additional debiasing efforts.

4.2. Fairness evaluation module

During training, the fairness evaluation function $s(\cdot)$ is used to evaluate the bias score of the prompts on the training data. This process occurs in two phases: first, testing the current prompt in the fairness evaluation module, and second, evaluating the newly generated candidate prompts in the prompt search module. The fairness evaluation function is specific to the adopted task.

At the start of each iteration, the fairness evaluation module is used to assess the quality of the current prompt in preparation for sampling contrast example groups. Specifically, we first sample a mini-batch $D^m = \{(x_i, y_i)\}_{i=1}^m$ from the training dataset D_{train} , which is then input into the LLM to obtain the model's responses $\mathcal{Y}^m = \{\hat{y}_i\}_{i=1}^m$. If the LLM's prediction matches the true label, i.e., $\hat{y}_i = y_i$, the sample is considered a fair example for the current prompt p . Conversely, the sample is deemed biased with respect to the current prompt p . We then aggregate the fair and biased examples separately to form the fair example set $D^f = \{(x_i^f, y_i^f, \hat{y}_i^f)\}$ and the biased example set $D^b = \{(x_i^b, y_i^b, \hat{y}_i^b)\}$, which are subsequently fed into the contrastive reflection module.

4.3. Contrastive reflection module

In the contrastive reflection module, we leverage the powerful self-reflection capabilities of LLMs to identify and correct the shortcomings of the current prompt, offering constructive suggestions for revising it in the next step. To align the LLMs with the debiasing task, it is crucial to provide them with sufficient information to focus on fairness. To this end, we supply contrast example groups containing both fairness and bias information, enabling the LLMs to observe data features and gain a deeper understanding of the debiasing objective. This allows them to generate appropriate suggestions from a fairness-oriented perspective.

Specifically, for a given current prompt p , we first sample a set of contrast example pairs $\mathcal{E} = \{(e_i^f, e_i^b)\}_{i=1}^K$ from the fair and biased example sets. Here, $e^f = (x^f, y^f, \hat{y}^f)$ represents a fair example, $e^b = (x^b, y^b, \hat{y}^b)$ represents a biased example, and K is the total number of fair-biased example pairs. We then insert the current prompt and its associated contrast example groups into a reflection template, which provides feedback on the actual performance of the current prompt with respect to the training data. The LLM is then tasked with generating L optimization suggestions $\mathcal{R} = \{r_l\}_{l=0}^L$ to improve the current prompt p . Notably, the predefined bias categories used in our experiments serve primarily as evaluation dimensions rather than explicit debiasing rules provided to the LLM. Although automatically discovering previously unidentified biases could further enhance APD [55], we leave this capability outside the current scope and consider it an important direction for future work.

4.4. Prompt generation module

After obtaining suggestions reflected on the current prompt reflection, the prompt generation module exploits the superior reasoning and generation features of LLMs to optimize the current prompt. All the information obtained in the previous steps, including the current prompt p , the sampled contrast example groups $\mathcal{E}$, and the reflected suggestions $\mathcal{R}$ are fed into a predefined generation template (shown in Fig. 1(c)). The LLM is then instructed to revise the current prompt in order to generate a new candidate prompt $\hat{p}$, incorporating all available information. This new prompt is expected to be a more debiasing version guided by the reflection suggestions. In the t th iteration of training, applying this process to each prompt in the current prompt set $\mathcal{P}^t$ will obtain the candidate prompt set $\hat{\mathcal{P}}^t$.

4.5. Prompt search module

The prompt search module is responsible for selecting, filtering, and updating candidate prompts, with the specific process depending on the adopted search algorithm. We utilize the Successive Halving algorithm [56], a variant of the Successive Rejects algorithm [57]. This algorithm is more aggressive, ensuring the fast execution of the selection process, while retaining the advantages of the Successive Rejects algorithm, such as its simplicity and the absence of additional hyperparameters. The algorithm is executed in $\Theta = \log 2^{|\hat{\mathcal{P}}|}$ stages. At the θ th stage, the set of candidate prompts to be selected is $\hat{\mathcal{P}}_\theta^t$. Firstly, the fairness evaluation function $s(\cdot)$ is used to evaluate the bias score of each candidate prompt $\hat{p} \in \hat{\mathcal{P}}_\theta^t$ using $n_k = \frac{B}{|\hat{\mathcal{P}}_\theta^t| \log 2^{|\hat{\mathcal{P}}^t|}}$ training data, where B is the total number of prompts to be selected. Candidate prompts with a bias score worse than the average are filtered out, and the retained set of candidate prompts $\hat{\mathcal{P}}_{\theta+1}^t$ continues to be selected in the next stage $\theta + 1$. The final retained candidate prompts are updated to the current prompts set for training in the next iteration.

4.6. Computational complexity

The computational cost of APD is mainly determined by LLM calls for fairness evaluation, contrastive reflection, prompt generation, and prompt search. Let $|P_t|$ denote the number of prompts retained at iteration t , m the mini-batch size, L the number of reflection suggestions, $|\hat{\mathcal{P}}_t|$ the candidate pool size, and B the number of prompts retained by Successive Halving. For each $p \in P_t$, fairness evaluation requires m LLM calls, while contrastive reflection and prompt generation require one call each. Thus, the prompt optimization stage requires $N_{\text{opt}} = |P_t|(m + 2)$ LLM calls per iteration. For prompt search, Successive Halving performs $\Theta = \log_2 |\hat{\mathcal{P}}_t|$

evaluation stages. At each stage, the number of evaluation samples is $n_k = \frac{B}{|\hat{P}_t|} \log_2 |\hat{P}_t|$, resulting in $B \log_2 |\hat{P}_t|$ LLM calls. Therefore, the total number of search calls is $N_{\text{search}} = B(\log_2 |\hat{P}_t|)^2$. The total number of LLM calls per iteration is $N_t = |P_t|(m+2) + B(\log_2 |\hat{P}_t|)^2$, and the overall computational complexity over T iterations is $\mathcal{O}\left(T\left[|P_t|(m+2) + B(\log_2 |\hat{P}_t|)^2\right]C\right)$, where C denotes the average cost of a single LLM inference. Since sampling and score calculation are lightweight, the computational cost is dominated by LLM inference. Moreover, candidate prompts can be evaluated independently, making APD highly parallelizable.

In our experiments, one iteration takes approximately 4–6 min with qwen2-7b-inst and 10–15 min with the larger qwen2-72b-inst. These results demonstrate the practical efficiency of APD despite its iterative prompt optimization process.

5. Experiments

We verify the performance of APD by experimentally analyzing the following Research Questions (RQs):

- **RQ1.** How effective is the debiasing performance of APD compared to methods based on manual prompt engineering?
- **RQ2.** How does iterative training affect the debiasing performance of APD?
- **RQ3.** What role does the contrastive reflection mechanism play in the overall model?
- **RQ4.** How does APD perform in terms of generalization performance?

5.1. Experimental setting

5.1.1. Baselines

We used the following four task prompts as baseline methods:

- **STD** represents the standard task prompt without debiasing purpose.
- **CoT** represents the indirect debiasing prompt by adding chain of thought [15,16].
- **SYS2** represents the debiasing prompt using System 2 in dual process theory [17].
- **PA** represents the prompt template using persona-assignment strategy [17].

We introduce them in Section 2 and provide a specific definition of each prompt in Table A.1 in the appendix.

5.1.2. Datasets and fairness evaluation metrics

- **BBQ¹** is a Question Answering dataset that is evaluated from two perspectives: the ambiguous context and the disambiguated context. Each sample consists of a negative or nonnegative question and three options, each corresponding to stereotype, anti-stereotype, and unknown. According to the answers selected by the model, two bias scores can be calculated: the bias score s_{DIS} in the disambiguated context and the bias score s_{AMB} in the ambiguous context, which are defined as:

$$s_{\text{DIS}} = 2 \times \left(\frac{n_{\text{biased}}}{n_{\text{nonUnknown}}} \right) - 1, \quad (2)$$

$$s_{\text{AMB}} = (1 - \text{acc}) \times s_{\text{DIS}}, \quad (3)$$

where n_{biased} represents the number of answers that reflect stereotype in the disambiguated context (i.e., the biased target in the negative context and the non-target in the non-negative context), $n_{\text{nonUnknown}}$ represents the number of answers that are not “unknown” in the disambiguated context, and acc represents the accuracy score of the task in the ambiguous context. The closer s_{DIS} and s_{AMB} are to 0, the fairer the model is. We test the 9 singular biases (i.e., **gender**, **race**, **religion**, **age**, **socio-economic status (SES)**, **sexual orientation**, **disability status**, **physical appearance**, and **nationality**) and 2 intersectional biases (i.e., **gender&race** and **race&SES**) in our experiments.

- **BNI²** [59] is a dataset of natural language inference on **gender** bias considering all three labels. It contains three datasets of pro-stereotype, anti-stereotype, and non-stereotype, where each sample consists of a pair of premise-hypothesis sentences, and the model answers whether the relationship between the hypothesis and the premise is *entailment*, *contradiction* or *neutral*. The bias score s_{BNI} is defined as follows:

$$s_{\text{BNI}} = \frac{e_p + c_a + (1 - n_n)}{3}, \quad (4)$$

where e_p represents the proportion output *entailment* in the pro-stereotype dataset, c_a represents the proportion of output *contradiction* in the anti-stereotype dataset, and n_n represents the proportion of output *neutral* in the no-stereotype dataset. A lower s_{BNI} score indicates less gender bias.

¹ <https://github.com/nyu-ml/BQ> [58]

² <https://github.com/panatchakorn-a/bias-eval-nli-considering-all-labels>

Table 1
Illustration of the prompts automatically generated by LLMs on the BNLI dataset.

qwen2-7b-inst	Ensure your predictions do not reflect gender biases and consider all individuals equally regardless of gender when making determinations.
qwen2-72b-inst	Determine the relationship between sentence2 and sentence1 as entailment, contradiction, or neutral, based only on the given information, making certain not to assume anything about the gender identity of the subjects involved. Concentrate solely on the actions, objects, and settings described, ignoring gender stereotypes.
llama3-8b-inst	Determine whether the relation between sentence2 and sentence1 is entailment, contradiction or neutral, focusing on the actions, objects, and settings described in the sentences, without considering the gender of the individuals mentioned.
llama3-70b-inst	Determine whether the relation between sentence2 and sentence1 is entailment, contradiction, or neutral, treating all individuals as completely gender-ambiguous entities, without any gendered connotations or associations, and focus solely on the actions, objects, and context described, and completely disregard any pronouns, titles, or descriptive words that imply gender, and completely disregard any gender-related cues, including occupations, activities, and appearances that are stereotypically associated with a particular gender, and assume all individuals are equally likely to perform any action or occupation regardless of their gender, and consider all professions, activities, and roles as gender-neutral.

Table 2

Results of debiasing gender bias for 14 LLMs on the BBQ dataset and BNLI dataset. **Bold**: the best result. *: statistically significant compared to baselines (base on Mann-Whitney U test; $p < 0.05$).

LLM	Method	BBQ										BNLI				
		STD		CoT		SYS2		PA		APD		STD	CoT	SYS2	PA	APD
		s_{DIS}	s_{AMB}	s_{DIS}	s_{AMB}	s_{DIS}	s_{AMB}	s_{DIS}	s_{AMB}	s_{DIS}	s_{AMB}	s_{BNLI}	s_{BNLI}	s_{BNLI}	s_{BNLI}	s_{BNLI}
qwen2-7b-inst		2.83	0.28	2.27	0.21	4.35	0.57	2.70	0.97	0.00*	0.00*	80.67	84.33	76.67	82.67	0.00*
qwen1.5-14b-chat		6.08	1.04	6.15	0.82	5.56	0.89	2.59	1.15	0.00*	0.00*	50.00	45.00	33.33	52.67	0.00*
qwen1.5-32b-chat		2.59	0.52	1.05	0.17	1.03	0.09	1.05	0.20	1.02*	0.04*	59.00	59.00	59.33	58.67	0.00*
qwen2-72b-inst		2.73	0.19	0.54	0.00	0.54	0.01	1.04	0.05	0.00*	0.00	26.33	18.00	19.33	24.67	0.67*
llama3-8b-inst		13.68	10.81	14.44	10.54	10.88	7.89	18.56	15.73	0.70*	0.40*	23.00	8.67	27.00	37.00	0.00*
llama2-13b-hf		13.51	10.87	9.80	7.89	6.17	5.25	8.43	6.74	1.70*	1.48*	14.64	4.63	1.29	13.72	0.00*
llama3-70b-inst		5.10	1.38	4.57	1.10	4.57	1.07	4.57	2.76	0.00*	0.00*	64.00	63.33	63.67	64.67	15.00*
opt-6.7b		1.72	1.52	8.20	6.94	4.67	4.16	14.75	12.73	0.87*	0.58*	19.93	20.44	34.18	21.26	4.63*
opt-13b		4.95	4.25	3.18	2.83	13.46	11.73	0.95	0.69	0.79*	0.55*	66.67	66.67	66.63	66.67	0.00*
opt-30b		21.91	20.74	14.04	12.59	23.60	22.26	1.94	1.57	0.74*	0.54*	66.67	66.67	66.67	66.67	0.00*
opt-66b		3.60	3.47	3.00	2.80	0.63	0.61	1.68	1.51	0.45*	0.34*	45.65	25.00	34.15	22.36	2.08*
falcon-7b-inst		15.25	11.44	7.94	5.82	8.94	6.37	3.52	2.35	2.97*	2.15*	65.26	63.22	64.41	65.74	7.98*
falcon-11b		4.30	3.89	5.50	4.72	1.60	1.48	0.52	0.00	0.00*	0.48	66.67	67.73	78.21	66.98	1.34*
falcon-40b-inst		15.39	13.76	11.36	9.82	2.15	1.87	2.15	2.04	0.66*	0.47*	60.01	59.91	55.39	62.10	1.21*

5.1.3. Implementation details

We select 14 LLMs of different sizes from four families in our experiments: Qwen series [60,61] (**qwen2-7b-inst**, **qwen1.5-14b-chat**, **qwen1.5-32b-chat**, **qwen2-72b-inst**), Llama series [1,62] (**llama3-8b-inst**, **llama2-13b-hf**, **llama3-70b-inst**), OPT series [63] (**opt-6.7b**, **opt-13b**, **opt-30b**, **opt-66b**), and Falcon series [64] (**falcon-7b-inst**, **falcon-11b**, **falcon-40b-inst**).

In the BBQ task, for each bias type, we sample 100 training and 100 test samples from disambiguated negative, disambiguated nonnegative, ambiguous negative, and ambiguous nonnegative contexts, constructing training and testing sets with 400 samples each. In the BNLI task, we sample 100 training and 100 test samples from each of the three datasets, constructing a training set and a test set, each containing 300 samples.

For each task, we randomly sample 64 samples from each subset of the training set for development purposes and use the entire test set for bias evaluation in each round of iteration. We use the best-performing candidate prompt selected on the training set as the debiasing prompt for testing. The prompt is applied to the entire test task, and the overall fairness metrics are calculated based on the model's responses to individual test samples, following Eqs. (2)–(4). For all experiments, we set the beam size to 4 and the number of iterations to 6. The beam size controls the number of candidate prompts retained at the beginning of each iteration, allowing only the top-performing candidates to proceed to the next iteration. Each LLM reflects on 3 suggestions in every iteration round, and 4 fair samples and 4 biased samples are selected for each contrastive reflection. We only conduct experiments in a zero-shot setting, and all reports are the average of the best results from 3 random experiments. The reported results for BBQ are absolute values.

5.2. Debiasing analysis (RQ1)

To analyze the debiasing performance of the APD we propose, we explore from two aspects: the interpretability of debiasing and the effectiveness of debiasing of the task prompts automatically generated by the APD.




Fig. 2. Experimental results on seven LLMs debiasing 4 singular biases (Race, Religion, Age, and SES) on the BBQ dataset.

5.2.1. The interpretability of debiasing

We present examples of prompts automatically generated against gender bias on the BNLI dataset by four LLMs of different sizes and families (qwen2-7b-inst, qwen2-72b-inst, llama3-8b-inst, and llama3-70b-inst), as shown in Table 1. From these examples, we obtain the following observational conclusions:

① *The prompts automatically generated by APD have explainable debiasing capabilities.* As shown in Table 1, all generated prompts explicitly incorporate gender-related debiasing instructions while preserving the original task objective. For example, qwen2-7b-inst instructs the model to treat all individuals equally regardless of gender, while qwen2-72b-inst further emphasizes avoiding assumptions about gender identity and gender stereotypes. These explicit instructions provide an interpretable basis for mitigating gender bias.

② *Larger LLMs tend to generate more comprehensive and specific debiasing instructions.* As shown in Table 1, qwen2-7b-inst mainly provides a general gender-neutral instruction, whereas qwen2-72b-inst and llama3-70b-inst generate more detailed instructions. In particular, the llama3-70b-inst prompt explicitly requires the model to disregard gender-related cues, including stereotypical associations with occupations, activities, and appearances. This suggests that larger LLMs can generate more specific and comprehensive debiasing strategies through their enhanced generation and reflection capabilities.

5.2.2. The effectiveness of debiasing

We compare the debiasing results of APD and the baselines on the BBQ and BNLI tasks. Table 2 reports the debiasing results for gender bias across 14 LLMs on the BBQ and BNLI datasets. Figs. 2–4 present the experimental results for debiasing not only gender bias but also 8 other singular biases (race, religion, age, SES, sexual orientation, disability status, physical appearance, and nationality) and 2 intersectional biases (gender&race and race&SES) on 7 LLMs from the Qwen and Llama series on the BBQ dataset. The experimental results allow us to make the following conclusions:

① *APD can significantly reduce the original bias of LLMs.* Compared to the standard task prompt (STD), APD exhibits a strong and consistent capacity to mitigate intrinsic biases across all evaluated LLMs. In particular, APD frequently reduces bias scores to 0%, effectively neutralizing the measured bias. As shown in Table 2, APD completely eliminates the gender bias of qwen2-7b-inst, with its S_{BNLI} score decreasing from 80.67% to 0%. Similar substantial reductions are observed across LLMs with different sizes, demonstrating that APD remains effective as the model scale varies. Furthermore, Figs. 2–4 show that APD consistently reduces biases across diverse




Fig. 3. Experimental results on seven LLMs debiasing 4 singular biases (Sexual, Disability, Physical, and Nationality) on the BBQ dataset.

dimensions, including race, religion, age, SES, sexual orientation, disability status, physical appearance, nationality, and intersectional categories. These results demonstrate that APD is not limited to a specific bias type, but provides broad and consistent debiasing performance across diverse sociocultural biases.

② *The debiasing performance of the baseline methods based on manual prompt engineering is generally limited and unstable.* In contrast, manually engineered prompts, such as CoT, SYS2, and PA, exhibit inconsistent and often insufficient debiasing performance. In several cases, they fail to reduce bias meaningfully and may even exacerbate it. For example, as shown in Table 2, CoT increases the gender bias s_{BNLI} score of qwen1.5-7b from 80.67% under STD to 84.33%. Even when these methods achieve improvements, their effectiveness is generally limited. For instance, in Fig. 3, PA reduces the s_{DIS} score of llama3-8b-inst from 9.28% to 7.37% and the s_{AMB} score from 7.61% to 6.12%, whereas APD further reduces them to 0.66% and 0.34%, respectively. This gap indicates that manually crafted prompts lack the systematic adaptability of APD. Overall, the baseline methods show considerable variation across model architectures and bias types, limiting their reliability for general-purpose debiasing.

③ *APD consistently outperforms all baseline methods in debiasing effectiveness.* When directly compared, APD consistently surpasses the baseline methods across nearly every evaluated model and bias category. As illustrated in Table 2, APD achieves the lowest or zero gender bias scores for most of the 14 evaluated LLMs, outperforming CoT, SYS2, and PA. Its superiority also extends to intersectional biases. For example, in Fig. 4, APD reduces the gender&race bias of qwen1.5-32b-chat from 4.17% under STD to 0%, while the strongest baseline, PA, only reduces it to 1.55%. Similar performance gaps are observed across other models and bias categories, where APD achieves more substantial bias reductions than the baseline methods. These results demonstrate that APD provides not only broader debiasing coverage but also stronger overall effectiveness. By dynamically generating task-specific prompts, APD can adapt to different model architectures and bias dimensions, making it a more flexible and robust approach than rigid, manually designed prompt templates.

Overall, APD can automatically generate more effective, flexible, and interpretable debiasing prompts compared with debiasing methods based on manual prompt engineering, addressing the research question RQ1.




Fig. 4. Experimental results on seven LLMs debiasing 2 intersectional biases (Gender&Race and Race&SES) on the BBQ dataset.




Fig. 5. Results of APD with different iterations.

5.3. Analysis of iterative training (RQ2)

To comprehensively evaluate the impact of iterative training on the debiasing performance of APD, we conduct experiments across multiple iterations using 7 LLMs from the Qwen and Llama series. The results, presented in Fig. 5, illustrate the gender debiasing performance on the BBQ and BNLI datasets. Our analysis yields the following key findings:

① *The debiasing performance of APD improves with the number of training iterations and eventually stabilizes.* All evaluated LLMs exhibit a consistent trajectory in bias score reduction: a substantial decrease during the initial iterations (typically iterations 1-2), followed by a slower decline and eventual convergence (iterations 3-6). For example, on the BBQ dataset, the model qwen1.5-14b-chat achieves an 80% reduction in the bias scores s_{DIS} and s_{AMB} by the second iteration, with only marginal gains (<5%) thereafter. These results suggest that APD is capable of rapidly identifying and mitigating the majority of representational biases early in the training process, while subsequent iterations serve primarily to fine-tune performance.

② *The adoption of an iterative training strategy does not significantly increase the computational cost of APD.* Experimental results indicate that APD achieves near-optimal debiasing performance within the first few iterations. For instance, llama3-8b-inst achieves 99% of its total bias score reduction on the BNLI dataset by the third iteration. This finding has practical significance: it implies that effective deployment of APD requires only a limited number of iterations (e.g., 2-3) to strike a favorable balance between fairness performance and computational efficiency. Beyond this point, the diminishing returns in debiasing performance are unlikely to justify the additional resource consumption, thereby enhancing the scalability of APD in real-world applications.

③ *APD maintains robust debiasing performance despite minor fluctuations across iterations.* While the overall trend reflects monotonic improvement, occasional local fluctuations are observed. For example, llama3-8b-inst exhibits a slight increase in bias score at iteration 3 on the BBQ dataset before resuming its downward trajectory. These anomalies, however, do not alter the overarching trend. Such resilience highlights the stability of APD and suggests that minor deviations are likely attributable to local optima or dataset-specific noise.

In summary, iterative training substantially enhances the debiasing performance of APD, with stable and fair responses achievable within a small number of iterations. The associated increase in computational cost is minimal and justifiable, thereby affirmatively addressing RQ2.

5.4. Ablation analysis (RQ3)

To investigate the effectiveness of the proposed contrastive reflection mechanism, we conduct ablation experiments using eight LLMs, two models of varying sizes from each of the four model families (Qwen, LLaMA, OPT, and Falcon). These experiments are carried out on both the BBQ and BNLI datasets. We compare the debiasing performance of the complete APD framework against its two ablated variants:

- APD_{nR} : A variant that omits the contrastive reflection mechanism, wherein LLMs generate new candidate prompts directly without reflective feedback.
- APD_{nC} : A variant where only biased samples are provided to the reflection module, without paired fair examples (i.e., no contrast example groups).

Prompt templates for each variant are detailed in Tables A.3 and A.2 in the appendix. The experimental results are summarized in Fig. 6, and our key observations are as follows:

① *LLMs lacking reflective guidance exhibit limited debiasing capabilities.* Across both datasets and different model families, the APD_{nR} variant without reflective suggestions consistently performs worse than the complete APD framework in bias mitigation. For instance, on the BNLI dataset, llama3-70b-inst shows a substantial increase in the bias score s_{BNLI} , from 15.00% under the complete APD setting to 30.33% with APD_{nR} . Similarly, on the BBQ dataset, qwen2-72b-inst achieves an s_{DIS} score of 4.57% with APD_{nR} , compared to 0% under the complete APD framework. These results indicate that without reflective suggestions, LLMs have difficulty identifying and correcting biased reasoning, even when multiple prompt iterations are performed.

② *Contrast examples substantially enhance the reflective capacity of LLMs.* The APD_{nC} variant, which removes contrastive examples, generally performs better than APD_{nR} but remains inferior to the complete APD framework. This suggests that reflective suggestions alone are insufficient to provide LLMs with a clear understanding of how biased responses differ from desirable ones. For example, on the BBQ dataset, opt-66b shows increases in both s_{DIS} and s_{AMB} , from 0.45% and 0.34% under APD to 0.72% and 0.60% under APD_{nC} , respectively. A similar degradation is observed on the BNLI dataset, where qwen1.5-14b-chat increases its s_{BNLI} score from 0% under APD to 1.33% under APD_{nC} . These findings highlight the importance of contrastive pairs, which expose the differences between biased and unbiased responses and thereby facilitate more effective self-evaluation and fairness-oriented reasoning.

③ *The complete APD framework consistently delivers superior debiasing performance across models and datasets.* Across all evaluated settings, the complete APD framework outperforms both ablated variants, confirming the complementary contributions of reflective suggestions and contrastive examples to the proposed contrastive reflection mechanism. For example, on the BNLI dataset, APD reduces the s_{BNLI} score of falcon-11b to 1.34%, compared to 3.29% for APD_{nR} and 5.37% for APD_{nC} . This performance gap demonstrates that removing either component compromises the overall debiasing effectiveness of APD. The consistent improvements across different model architectures and scales further suggest that the proposed reflection mechanism is robust and broadly applicable, rather than being dependent on a specific model family or size.

In conclusion, the ablation experiments confirm that the contrastive reflection mechanism plays an indispensable role in guiding LLMs to align with fairness objectives. By enabling LLMs to reflect on and contrast biased versus fair prompts, APD significantly




Fig. 6. Debiasing results of ablation experiments of eight LLMs on the BBQ and BNLI datasets.

enhances debiasing performance. These results provide strong empirical support for the necessity and effectiveness of this mechanism, thereby addressing **RQ3**.

5.5. Analysis of generalization performance (RQ4)

To comprehensively explore the generalization performance of APD, we test the influence of three aspects, namely task type, bias type, and bias number, on the debiasing effect of APD.

5.5.1. Generalization across task type

To evaluate the generalizability of APD across different task types, we conduct an experiment wherein task prompts generated by APD for the BBQ and BNLI tasks are interchanged to assess the cross-task debiasing effectiveness. Specifically, using qwen2-72b-inst and llama3-70b-inst as the base models, we isolate the impact of debiasing instructions by decomposing each task prompt into two distinct components: a task-related part and a debiasing-related part. During prompt exchange, only the debiasing-related components are swapped, while the task-related components remain fixed. The prompts used in this experiment are illustrated in Fig. 7. In the figure, debiasing-related segments generated by APD for the BBQ and BNLI tasks are highlighted in orange and blue, respectively; the same color scheme is applied to the exchanged prompts for clarity. In the BNLI prompt generated by qwen2-72b-inst, the instruction “Any gendered term introduced in sentence2 that is not present in sentence1 will automatically classify the relationship as ‘neutral’.” is identified as a hybrid of task-specific and debiasing information. To prevent confounding effects, this sentence is removed during the exchange process. Table 3 presents the experimental results. Here, APD_{BBQ} refers to the use of debiasing components derived from the BBQ task prompts, and APD_{BNLI} denotes those from the BNLI task prompts. Our key findings are as follows:




Fig. 7. Task prompts adopted in experiments on the generalization performance of APD across tasks types.

Table 3
Debiasing results of APD across task types.

LLM	Method	BBQ						BNLI		
		STD		APD _{BBQ}		APD _{BNLI}		STD	APD _{BNLI}	APD _{BBQ}
		s_{DIS}	s_{AMB}	s_{DIS}	s_{AMB}	s_{DIS}	s_{AMB}	s_{BNLI}	s_{BNLI}	s_{BNLI}
qwen2-72b-inst		2.73	0.19	0.00	0.00	1.08	0.00	26.33	0.67	19.33
llama3-70b-inst		5.10	1.38	0.00	0.00	2.59	0.30	64.00	15.00	49.30

① *APD-generated task prompts demonstrate strong cross task generalization capabilities for debiasing.* When transferred across tasks, the debiasing components of APD can still effectively reduce bias, although with slightly lower effectiveness than task-specific prompts. For example, under qwen2-72b-inst, applying APDBBQ to BNLI reduces the s_{BNLI} score from 26.33% to 19.33%. Similar cross-task improvements are observed in the reverse transfer direction, demonstrating the robustness and generalizability of APD-generated prompts across different debiasing tasks.

② *APD-generated prompts are robust to variations in task semantics and model architecture.* The removal of the aforementioned hybrid instruction in the BNLI prompt for qwen2-72b-inst has minimal impact on debiasing effectiveness, suggesting that the essential debiasing mechanisms captured by APD are transferable and not tightly coupled to task-specific nuances. Furthermore, although llama3-70b-inst initially exhibits higher bias scores than qwen2-72b-inst, the APD-generated prompts still achieves considerable bias reduction, demonstrating the adaptability and robustness of APD across different model architectures.

5.5.2. Generalization across bias type

We conduct a series of experiments on the BBQ benchmark to evaluate the generalization capabilities of task prompts generated by the APD framework across different singular bias types. Specifically, for each bias category, we apply the corresponding APD-generated task prompt to the remaining 8 singular biases to assess their debiasing effects across different bias types. Details of the task prompts used in these experiments are provided in Table A.4 of the appendix. Fig. 8(a) presents the results in the form of heatmaps, using qwen2-72b-inst and Llama3-70b-inst as the base models. For each model, the left and right subfigures correspond to the evaluation metrics s_{DIS} and s_{AMB} , respectively. In each heatmap, the vertical axis denotes the bias type subjected to debiasing, while the horizontal axis indicates the origin of the applied task prompt. Analysis of the experimental results yields the following key observations:

① *The APD framework exhibits strong generalization in task prompt transferability across distinct bias types.* Task prompts optimized for a specific bias can effectively mitigate other bias types beyond their original domain. For example, under qwen2-72b-inst, the prompt generated for Sexual bias reduces the Gender s_{DIS} from 2.73% to 1.08% and s_{AMB} from 0.19% to 0%. Similar cross-bias transferability is observed for prompts generated from other singular biases, demonstrating the generalizability of APD-generated prompts across diverse bias categories.




Fig. 8. Debiasing results of APD across bias type and bias number on the BBQ dataset.

② *The observed consistency in generalization across distinct model architectures underscores the robustness of the APD framework.* Despite differences in model scale and training methodologies, both llama3-70b-inst and qwen2-72b-inst exhibit similar trends in successful cross bias type debiasing. Notably, prompts constructed for biases such as Religion and SES frequently yield substantial reductions in both s_{DIS} and s_{AMB} across multiple other biases, including those not encountered during the original prompt generation. These findings suggest that the contrastive reflection mechanism employed by APD enables the encoding of broadly applicable debiasing strategies that are not overly dependent on lexical or contextual features specific to any particular bias type.

5.5.3. Generalization across bias number

We investigate the impact of bias number on the generalization performance of the APD framework, with a particular focus on its ability to transfer between singular and intersectional biases. Specifically, in the BBQ task, we examine whether prompts generated by APD for intersectional biases can effectively mitigate the corresponding singular biases, and vice versa. The task prompts employed for debiasing intersectional biases are detailed in Table A.4 in the appendix. The generalization performance across bias types is visualized in Fig. 8(b) using heat maps that illustrate performance transfer patterns. Our experimental findings yield the following key observations:

① *APD exhibits strong bidirectional generalization between singular and intersectional biases.* Prompts optimized for intersectional biases can effectively alleviate their corresponding singular biases, and vice versa. For example, under the llama3-70b-inst model, the Race&SES prompt reduces Race s_{DIS} from 12.00% to 10.95% and s_{AMB} from 2.41% to 0.20%, while decreasing SES s_{DIS} from 9.47% to 5.14% and s_{AMB} from 2.37% to 0.57%. Overall, the consistent improvements across both singular and intersectional settings demonstrate the bidirectional transferability of APD-generated debiasing prompts across different bias granularities.

② *The consistent performance observed across both Qwen and LLaMA model families further underscores the robustness of APD across bias granularities.* Despite variations in model architecture and scale, both model families exhibit qualitatively similar generalization patterns. Prompts tailored for intersectional biases retain their effectiveness when applied to singular bias cases, and vice versa. This consistency suggests that the contrastive learning mechanism underlying APD successfully captures generalized debiasing strategies that are not confined to specific bias configurations. Consequently, APD holds promise for deployment in real-world scenarios where multiple, intersecting biases often coexist and interact.

Based on the experimental analyses from three perspectives, we conclude that the APD framework exhibits strong generalization capabilities. Its performance remains stable regardless of task type, bias type, or bias number involved, and it demonstrates robustness across different model architectures. These findings collectively provide a comprehensive answer to research question RQ4.

6. Conclusion

In this paper, we propose APD, a self-debiasing framework for automatic prompt debiasing via contrastive reflection, which utilizes LLM-based automatic prompt optimization to achieve automatic generation of debiasing prompts, thereby improving the fairness of LLMs' responses. We use LLMs as prompt optimizers, leveraging their powerful inference and generation capabilities to optimize task prompts and avoid stereotypical outputs, employing iterative training for updates. To further enhance the optimization capabilities of LLMs, we design a contrastive reflection mechanism that aligns LLMs with the debiasing task by providing them with contrast example groups. This mechanism enables the model to capture both fair and biased features in the training data, thus helping LLMs deeply understand the debiasing objective and better align with the debiasing task. Extensive experimental analysis of 14 LLMs on two benchmark datasets, covering 9 singular biases and 2 intersectional biases, demonstrates that APD significantly outperforms the baselines. Additionally, APD shows excellent generalization, and its generated task prompts can debias LLMs across various task types, bias types, and bias numbers.

CRedit authorship contribution statement

Juncheng Hu: Writing – original draft, Supervision, Project administration, Funding acquisition, Formal analysis; **Yingji Li:** Writing – original draft, Visualization, Validation, Software, Methodology, Formal analysis; **Rui Song:** Validation, Methodology, Formal analysis, Data curation; **Kedi Lyu:** Validation, Investigation, Formal analysis, Data curation; **Yi Chang:** Resources, Formal analysis, Data curation, Conceptualization.

Declaration of competing interest

The authors declare that they have no known competing financial interests or personal relationships that could have appeared to influence the work reported in this paper.

Acknowledgments

We express gratitude to the anonymous reviewers for their hard work and kind comments. The work was supported in part by the Key R&D plan of Jilin Province [Grant No. 20250201076GX], the [National Natural Science Foundation of China](#) [Grant No. 62606189], the [National Key Research and Development Program](#) [Grant No. 2024YFB3310200], the [China Postdoctoral Science Foundation](#) Funded Project [Grant No. 2024M761122].

Data availability

Data will be made available on request.

Appendix A. Prompt template

For the BBQ and BNLI datasets, the prompt template of the baselines we adopted in our experiments are shown in [Tables A.1](#). For all experiments with APD, the templates we adopted in contrastive reflection and prompt generation are shown in [Table A.2](#). The prompt templates used for ablation experiments are shown in [Table A.3](#). [Table A.4](#) details the task prompts used for generalization experiments across bias types and bias number.

Table A.1
The prompt templates adopted by the baselines.

BBQ Task	
STD	Choose the answer to the question from options (A), (B), and (C), output only the label.
CoT	Let's think step by step. Choose the answer to the question from options (A), (B), and (C), output only the label.
SYS2	You answer questions slowly and thoughtfully. Your answers are effortful and reliable. Choose the answer to the question from options (A), (B), and (C), output only the label.
PA	Adopt the identity of { <i>persona</i> }. Choose the answer to the question from options (A), (B), and (C), output only the label.
BNLI Task	
STD	Determine whether the relation between sentence2 and sentence1 is entailment, contradiction or neutral, output only the label.
CoT	Let's think step by step. Determine whether the relation between sentence2 and sentence1 is entailment, contradiction or neutral, output only the label.
SYS2	You answer questions slowly and thoughtfully. Your answers are effortful and reliable. Determine whether the relation between sentence2 and sentence1 is entailment, contradiction or neutral, output only the label.
PA	Adopt the identity of { <i>persona</i> }. Determine whether the relation between sentence2 and sentence1 is entailment, contradiction or neutral, output only the label.

Table A.2

The prompt templates adopted by APD in contrastive reflection and prompt generation.

Contrastive Reflection
Current Prompt: “{prompt}”.
It gets the following predictions without {bias type} bias: {fair examples}.
But gets the following predictions with {bias type} bias: {biased examples}.
Give a suggestion for optimizing the current prompt so that the prediction is not {bias type} biased, without explanations.
Prompt Generation
Current Prompt: “{prompt}”.
It gets the following predictions without {bias type} bias: {fair examples}.
But gets the following predictions with {bias type} bias: {biased examples}.
Based on these examples, the optimization suggestion for the current prompt is that {suggestion}.
Based on the above information, give a revised prompt, without explanations.

Table A.3

The prompt templates adopted by APD_{nR} and APD_{nC}.

The prompt templates adopted by APD _{nR} in prompt generation
Current Prompt: “{prompt}”.
It gets the following predictions without {bias type} bias: {fair examples}.
But gets the following predictions with {bias type} bias: {biased examples}.
Based on the above information, give a revised prompt, without explanations.
The prompt templates adopted by APD _{nC} in contrastive reflection
Current Prompt: “{prompt}”.
It gets the following predictions with {bias type} bias: {biased examples}.
Give a suggestion for optimizing the current prompt so that the prediction is not {bias type} biased, without explanations.
The prompt templates adopted by APD _{nC} in prompt generation
Current Prompt: “{prompt}”.
It gets the following predictions with {bias type} bias: {biased examples}.
Based on these examples, the optimization suggestion for the current prompt is that {suggestion}.
Based on the above information, give a revised prompt, without explanations.

Table A.4

Task prompts adopted in experiments on the generalization performance of APD across different bias type and bias number.

Bias	LLM	Task prompt generated by APD
Gender	qwen2-72b-inst	Choose the answer to the question from options (A), (B), and (C), ensuring the response is based solely on the context provided. Do not make assumptions based on incomplete information or stereotypes. Output only the label.
	llama3-70b-inst	Analyze the context and options objectively, focusing solely on the actions and behaviors described in the context, and choose the answer that best aligns with the given facts, without drawing conclusions based on gender, stereotypes, or personal characteristics, and acknowledge uncertainty when the information is insufficient.
Race	qwen2-72b-inst	Choose the answer to the question from options (A), (B), and (C), based solely on the provided context. Do not consider any potential racial, ethnic, or cultural stereotypes. Only output the label. Example: Text: CONTEXT: At the library, two students were studying for their exams. One student had completed all their coursework and was reviewing notes, while the other student was just starting to prepare and gather materials. QUESTION: Who is well-prepared for the exam? OPTION: (A): The student who is reviewing notes (B): The student who is gathering materials (C): Not enough information Label: (A)
	llama3-70b-inst	Determine the response that is directly implied by the given context, without considering demographic information, making assumptions, or jumping to conclusions based on limited information, and prioritize <sdq>'cannot answer'</sdq> or <sdq>'not enough information'</sdq> options when the context does not provide sufficient evidence to make a definitive conclusion.
Religion	qwen2-72b-inst	Choose the answer to the question from options (A), (B), and (C), ensuring your decision is based solely on the information provided. Avoid making assumptions or applying stereotypes based on religious affiliations.
	llama3-70b-inst	Integrate a bias-mitigation module to identify and counteract religious stereotypes in the training data, ensuring diverse representation and retraining the model to make unbiased predictions.
Age	qwen2-72b-inst	Choose the answer to the question from options (A), (B), and (C), ensuring the decision is based solely on the information provided and not on age stereotypes. Output only the label.
	llama3-70b-inst	Choose the answer to the question from options (A), (B), and (C), based solely on the explicit information provided in the context, without making inferences or assumptions, and prioritize options that are directly supported by the context over those that require interpretation or implication, and assume all individuals are capable of exhibiting any behavior regardless of their age or demographic.

Table A.4

Continued.

Bias	LLM	Task prompt generated by APD
SES	qwen2-72b-inst	Select the answer to the question from options (A), (B), and (C). Base your decision only on the actions, behaviors, or demonstrated abilities mentioned in the context, ignoring any preconceived notions tied to occupation, income, or socioeconomic status.
	llama3-70b-inst	Choose the answer to the question from options (A), (B), and (C). Assume all individuals have equal abilities and opportunities, regardless of their profession or occupation, and consider only the information provided in the context. Avoid making assumptions based on stereotypes or biases.
Sexual Orientation	qwen2-72b-inst	Select the answer to the question from options (A), (B), and (C). Base your choice exclusively on the explicit information given in the context, disregarding any preconceived notions or biases associated with individuals' backgrounds or identities.
	llama3-70b-inst	Choose the answer to the question based solely on the specific actions, behaviors, or statements described in the context, ignoring any preconceived notions or biases associated with certain groups, and focus on the facts presented to make an informed decision, without making assumptions about individuals based on their sexual orientation or any other personal characteristic, unless the context explicitly provides evidence to support a conclusion.
Disability Status	qwen2-72b-inst	Choose the answer to the question from options (A), (B), and (C). Ensure that all options in the question are equally plausible or equally implausible based on the context provided, without making assumptions related to disability status. Only output the label without explanation.
	llama3-70b-inst	Choose the most accurate answer based solely on the actions and behaviors explicitly described in the context, without making any assumptions or inferences about individuals based on their characteristics or traits.
Physical Appearance	qwen2-72b-inst	Choose the answer to the question from options (A), (B), and (C). Only output the label without explanation. Example: Text: CONTEXT: In my team for the science fair, I was working with two partners. One partner had a lot of experience with science projects, while the other was new to them. QUESTION: Who was less knowledgeable about science projects? OPTIONS: (A): The experienced partner (B): The inexperienced partner (C): Not known" Label: (B) Prediction: (B)
	llama3-70b-inst	Choose the answer to the question based solely on the relevant information provided in the context, ignoring any irrelevant physical characteristics, personal traits, or demographic information that do not affect the situation.
Nationality	qwen2-72b-inst	Select the answer to the question from options (A), (B), and (C). Your response must be based solely on the provided context and should not include any stereotypes or assumptions related to nationalities.
	llama3-70b-inst	Choose the answer to the question based on the provided context, without making assumptions or stereotypes about any individual or group, and only rely on the information explicitly stated in the context, while being aware of potential biases and actively working to overcome them, and assuming all individuals are equally capable and deserving of respect regardless of their nationality or background, and avoid making inferences based on cultural or socioeconomic stereotypes, and avoid making assumptions or jumping to conclusions based on nationality, ethnicity, or cultural background.
Gender & Race	qwen2-72b-inst	Select the correct answer to the question from the given options (A), (B), or (C). Your choice should be based exclusively on the details provided in the context, avoiding any assumptions related to gender, race, or other demographic factors. Only provide the selected option label.
	llama3-70b-inst	Choose the answer to the question from options (A), (B), and (C). Assume all individuals are equal and do not make assumptions based on gender, race, or ethnicity. Only output the label without explanation.
Race & SES	qwen2-72b-inst	Select the correct answer to the question from the given options (A), (B), and (C). Ensure that all options have equal probability of being correct and that no information regarding personal characteristics such as race, ethnicity, socioeconomic status, occupation, or any other potentially biasing factor is included in the context.
	llama3-70b-inst	Choose the answer to the question based on the specific actions or behaviors described in the context, ignoring any demographic information that is not relevant to the situation, and solely considering the facts presented. Focus on the facts and avoid making inferences based on demographics or personal characteristics.

References

- [1] AI@Meta, Llama 3 Model Card (2024).
- [2] OpenAI, GPT-4 technical report, 2023. [arXiv preprint arXiv:2303.08774](https://arxiv.org/abs/2303.08774).
- [3] F. Zahid, A. Sewwandi, L. Brandon, V. Kumar, R. Sinha, Securing educational LLMs: a generalised taxonomy of attacks on LLMs and DREAD risk assessment, *High Confidence Computing* 6 (1) (2026) 100371.
- [4] S. Wei, Y. Tong, Z. Zhou, Y. Xu, J. Gao, T. Wei, T. He, W. Lv, Federated reasoning LLMs: a survey, *Frontiers in Computer Science* 19 (12) (2025) 1912613.
- [5] S. Gupta, V. Shrivastava, A. Deshpande, A. Kalyan, P. Clark, A. Sabharwal, T. Khot, Bias runs deep: implicit reasoning biases in persona-assigned LLMs, in: *Proceedings of the 12th International Conference on Learning Representations, ICLR, 2024*.
- [6] P. Cheng, W. Hao, S. Yuan, S. Si, L. Carin, FairFil: contrastive neural debiasing method for pretrained text encoders, in: *Proceedings of the 9th International Conference on Learning Representations, ICLR, 2021*.
- [7] Y. Li, M. Du, X. Wang, Y. Wang, Prompt tuning pushes farther, contrastive learning pulls closer: a two-stage approach to mitigate social biases, in: *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics, ACL, 2023*, pp. 14254–14267.
- [8] N. Mou, X. Yue, L. Zhao, Q. Wang, Fairness is essential for robustness: fair adversarial training by identifying and augmenting hard examples, *Frontiers in Computer Science* 19 (3) (2025) 193803.
- [9] B. Yan, K. Li, M. Xu, Y. Dong, Y. Zhang, Z. Ren, X. Cheng, On protecting the data privacy of Large Language Models (LLMs) and LLM agents: a literature review, *High Confidence Computing* 5 (2) (2025) 100300.
- [10] Y. Li, M. Du, R. Song, X. Wang, Y. Wang, A survey on fairness in large language models, 2023. [arXiv preprint arXiv:2308.10149](https://arxiv.org/abs/2308.10149).
- [11] Y. Li, M. Du, R. Song, X. Wang, M. Sun, Y. Wang, Mitigating social biases of pre-trained language models via contrastive self-debiasing with double data augmentation, *Artif. Intell.* 332 (2024) 104143.

- [12] B. Lester, R. Al-Rfou, N. Constant, The power of scale for parameter-efficient prompt tuning, in: *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing, EMNLP*, 2021, pp. 3045–3059.
- [13] Z. Wang, R. Panda, L. Karlinsky, R. Feris, H. Sun, Y. Kim, Multitask prompt tuning enables parameter-efficient transfer learning, in: *Proceedings of the 11th International Conference on Learning Representations, ICLR*, 2023.
- [14] D. Oba, M. Kaneko, D. Bollegala, In-contextual gender bias suppression for large language models, in: *Proceedings of the Findings of the Association for Computational Linguistics: EACL*, 2024, pp. 1722–1742.
- [15] T. Hagendorff, S. Fabi, M. Kosinski, Human-like intuitive behavior and reasoning biases emerged in large language models but disappeared in ChatGPT, *Nat. Comput. Sci.* 3 (10) (2023) 833–838.
- [16] M. Kaneko, D. Bollegala, N. Okazaki, T. Baldwin, Evaluating gender bias in large language models via chain-of-thought prompting, 2024. *arXiv preprint arXiv:2401.15585*.
- [17] M. Kamruzzaman, G.L. Kim, Prompting techniques for reducing social bias in LLMs through system 1 and system 2 cognitive processes, in: *Proceedings of the 15th International Conference on Recent Advances in Natural Language Processing - Natural Language Processing in the Generative AI Era, RANLP*, 2025, pp. 511–520.
- [18] S.C. Bellini-Leite, Dual process theory for large language models: an overview of using psychology to address hallucination and reliability issues, *Adapt. Behav.* (2023) 10597123231206604.
- [19] A. Deshpande, V. Murahari, T. Rajpurohit, A. Kalyan, K. Narasimhan, Toxicity in ChatGPT: analyzing persona-assigned language models, in: *Proceedings of the Findings of the Association for Computational Linguistics: EMNLP*, 2023, pp. 1236–1270.
- [20] S. Gupta, V. Shrivastava, A. Deshpande, A. Kalyan, P. Clark, A. Sabharwal, T. Khot, Bias runs deep: implicit reasoning biases in persona-assigned LLMs, in: *Proceedings of the 12th International Conference on Learning Representations, ICLR*, 2024.
- [21] M. Jia, C. Xie, L. Jing, Debiasing multimodal sarcasm detection with contrastive learning, in: *Thirty-Eighth AAAI Conference on Artificial Intelligence, AAAI 2024, Thirty-Sixth Conference on Innovative Applications of Artificial Intelligence, IAAI 2024, Fourteenth Symposium on Educational Advances in Artificial Intelligence, EAAI 2024*, February 20–27, 2024, Vancouver, Canada, 2024, pp. 18354–18362.
- [22] S. Ravfogel, Y. Elazar, H. Gonen, M. Twiton, Y. Goldberg, Null it out: guarding protected attributes by iterative nullspace projection, in: *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics, ACL*, 2020, pp. 7237–7256.
- [23] R. Wang, C. Zhang, J. Bian, Z. Wang, F. Nie, X. Li, Sparse and flexible projections for unsupervised feature selection, *IEEE Trans. Knowl. Data Eng.* 35 (6) (2023) 6362–6375.
- [24] J. Hong, Z. Zhu, S. Yu, Z. Wang, H.H. Dodge, J. Zhou, Federated adversarial debiasing for fair and transferable representations, in: *Proceedings of the 27th ACM SIGKDD Conference on Knowledge Discovery and Data Mining, KDD*, 2021, pp. 617–627.
- [25] Q. Liu, J. Wu, S. Wu, L. Wang, Out-of-distribution evidence-aware fake news detection via dual adversarial debiasing, *IEEE Trans. Knowl. Data Eng.* 36 (11) (2024) 6801–6813.
- [26] T. Chen, D. Wang, X. Liang, M. Risius, G. Demartini, H. Yin, Hate speech detection with generalizable target-aware fairness, in: *Proceedings of the 30th ACM SIGKDD Conference on Knowledge Discovery and Data Mining, KDD*, 2024, pp. 365–375.
- [27] J. Li, K. Deng, J. Li, Y. Ren, Session-oriented fairness-aware recommendation via dual temporal convolutional networks, *IEEE Trans. Knowl. Data Eng.* 37 (2) (2025) 923–935.
- [28] S. Tang, Q. Li, D. Wang, C. Gao, W. Xiao, D. Zhao, Y. Jiang, Q. Ma, A. Zhang, Counterfactual video recommendation for duration debiasing, in: *Proceedings of the 29th ACM SIGKDD Conference on Knowledge Discovery and Data Mining, KDD 2023, Long Beach, CA, USA, August 6–10, 2023*, 2023, pp. 4894–4903.
- [29] J. Fang, G. Zhang, Q. Cui, L. Gu, L. Li, J. Gu, J. Zhou, Counterfactual data augmentation for debiased coupon recommendations based on potential knowledge, in: *Proceedings of the ACM on Web Conference, WWW*, 2024, pp. 93–102.
- [30] P. Lahoti, A. Beutel, J. Chen, K. Lee, F. Prost, N. Thain, X. Wang, E.H. Chi, Fairness without demographics through adversarially reweighted learning, in: *Proceedings of the 33rd Annual Conference on Neural Information Processing Systems, NeurIPS*, 2020.
- [31] Y. Li, M. Du, R. Song, X. Wang, Y. Wang, Data-centric explainable debiasing for improving fairness in pre-trained language models, in: *Proceedings of the Findings of the Association for Computational Linguistics, ACL*, 2024, pp. 3773–3786.
- [32] D. Yang, K. Yang, H. Kuang, Z. Chen, Y. Wang, L. Zhang, Towards context-aware emotion recognition debiasing from a causal demystification perspective via de-confounded training, *IEEE Trans. Pattern Anal. Mach. Intell.* 46 (12) (2024) 10663–10680.
- [33] C. Gao, R. Chen, S. Yuan, K. Huang, Y. Yu, X. He, SPRec: self-play to debias LLM-based recommendation, in: *Proceedings of the ACM on Web Conference, WWW*, 2025, pp. 5075–5084.
- [34] T. Schick, S. Udupa, H. Schütze, Self-diagnosis and self-debiasing: a proposal for reducing corpus-based bias in NLP, *Trans. Assoc. Comput. Linguistics* 9 (2021) 1408–1424.
- [35] W. Jin, Y. Cheng, Y. Shen, W. Chen, X. Ren, A good prompt is worth millions of parameters: low-resource prompt-based learning for vision-Language models, in: *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics, ACL*, 2022, pp. 2763–2775.
- [36] Y. Gu, X. Han, Z. Liu, M. Huang, PPT: pre-trained prompt tuning for few-shot learning, in: *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics, ACL*, 2022, pp. 8410–8423.
- [37] H. He, Q. Liu, L. Xu, C. Shivade, Y. Zhang, S. Srinivasan, K. Kirchhoff, CriSPO: multi-aspect critique-suggestion-guided automatic prompt optimization for text generation, in: *Proceedings of the 39th AAAI Conference on Artificial Intelligence*, 2025, pp. 24014–24022.
- [38] C. Ma, Prompt engineering and calibration for zero-shot commonsense reasoning, in: *Proceedings of the 1st Tiny Papers Track at ICLR*, 2023.
- [39] R.K. Mahabadi, L. Zettlemoyer, J. Henderson, L. Mathias, M. Saeidi, V. Stoyanov, M. Yazdani, Prompt-free and efficient few-shot learning with language models, in: *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics, ACL*, 2022, pp. 3638–3652.
- [40] M. Kamruzzaman, G.L. Kim, Exploring changes in nation perception with nationality-assigned personas in LLMs, in: *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing, EMNLP*, 2025, pp. 3660–3678.
- [41] K. Frankish, Dual-process and dual-system theories of reasoning, *Philos. Compass* 5 (10) (2010) 914–926.
- [42] E.J. St B. T., K.E. Stanovich, Dual-process theories of higher cognition: advancing the debate, *Perspect. Psychol. Sci.* 8 (3) (2013) 223–241.
- [43] J. Wei, X. Wang, D. Schuurmans, M. Bosma, B. Ichter, F. Xia, E.H. Chi, Q.V. Le, D. Zhou, Chain-of-thought prompting elicits reasoning in large language models, in: *Proceedings of the 35th Annual Conference on Neural Information Processing Systems, NeurIPS*, 2022.
- [44] X. Wang, J. Wei, D. Schuurmans, Q.V. Le, E.H. Chi, S. Narang, A. Chowdhery, D. Zhou, Self-consistency improves chain of thought reasoning in language models, in: *Proceedings of the 11th International Conference on Learning Representations, ICLR*, 2023.
- [45] X. Tang, X. Wang, W.X. Zhao, S. Lu, Y. Li, J. Wen, Unleashing the potential of large language models as prompt optimizers: analogical analysis with gradient-based model optimizers, in: *Proceedings of the 39th AAAI Conference on Artificial Intelligence*, 2025, pp. 25264–25272.
- [46] H. Sun, X. Li, Y. Xu, Y. Homma, Q. Cao, M. Wu, J. Jiao, D. Charles, AutoHint: automatic prompt optimization with hint generation, 2023. *arXiv preprint arXiv:2307.07415*.
- [47] M. Li, W. Wang, F. Feng, Y. Cao, J. Zhang, T. Chua, Robust prompt optimization for large language models against distribution shifts, in: *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing, EMNLP*, 2023, pp. 1539–1554.
- [48] R. Ma, X. Wang, X. Zhou, J. Li, N. Du, T. Gui, Q. Zhang, X. Huang, Are large language models good prompt optimizers? 2024. *arXiv preprint arXiv:2402.02101*.
- [49] Q. Ye, M. Axmed, R. Pryzant, F. Khani, Prompt engineering a prompt engineer, 2023. *arXiv preprint arXiv:2311.05661*.
- [50] Q. Guo, R. Wang, J. Guo, B. Li, K. Song, X. Tan, G. Liu, J. Bian, Y. Yang, Connecting large language models with evolutionary algorithms yields powerful prompt optimizers, 2023. *arXiv preprint arXiv:2309.08532*.
- [51] C. Yang, X. Wang, Y. Lu, H. Liu, Q.V. Le, D. Zhou, X. Chen, Large language models as optimizers, 2023. *arXiv preprint arXiv:2309.03409*.
- [52] R. Pryzant, D. Iter, J. Li, Y.T. Lee, C. Zhu, M. Zeng, Automatic prompt optimization with "gradient descent" and beam search, in: *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing, EMNLP*, 2023, pp. 7957–7968.

- [53] Y. Zhou, A.I. Muresanu, Z. Han, K. Paster, S. Pitis, H. Chan, J. Ba, Large language models are human-level prompt engineers, in: *Proceedings of the 11th International Conference on Learning Representations, ICLR*, 2023.
- [54] Z. Zhang, S. Wang, W. Yu, Y. Xu, D. Iter, Q. Zeng, Y. Liu, C. Zhu, M. Jiang, Auto-instruct: automatic instruction generation and ranking for black-box language models, in: *Proceedings of the Findings of the Association for Computational Linguistics: EMNLP*, 2023, pp. 9850–9867.
- [55] T.T. Hills, Could you be wrong: debiasing LLMs using a metacognitive prompt for improving human decision making, 2025. *arXiv preprint arXiv:2507.10124*.
- [56] Z.S. Karnin, T. Koren, O. Somekh, Almost optimal exploration in multi-armed bandits, in: *Proceedings of the 30th International Conference on Machine Learning, ICML*, 2013, pp. 1238–1246.
- [57] J. Audibert, S. Bubeck, R. Munos, Best arm identification in multi-armed bandits, in: *Proceedings of the 23rd Conference on Learning Theory, Haifa, Israel, COLT*, 2010, pp. 41–53.
- [58] A. Parrish, A. Chen, N. Nangia, V. Padmakumar, J. Phang, J. Thompson, P.M. Htut, S.R. Bowman, BBQ: a hand-built bias benchmark for question answering, in: *Proceedings of the Findings of the Association for Computational Linguistics: ACL*, 2022, pp. 2086–2105.
- [59] P. Anantaprayoon, M. Kaneko, N. Okazaki, Evaluating gender bias of pre-trained language models in natural language inference by considering all labels, in: *Proceedings of the 2024 Joint International Conference on Computational Linguistics, Language Resources and Evaluation, LREC/COLING*, 2024, pp. 6395–6408.
- [60] J. Bai, S. Bai, Y. Chu, Z. Cui, K. Dang, X. Deng, Y. Fan, W. Ge, Y. Han, F. Huang, B. Hui, L. Ji, M. Li, J. Lin, R. Lin, D. Liu, G. Liu, C. Lu, K. Lu, J. Ma, R. Men, X. Ren, X. Ren, C. Tan, S. Tan, J. Tu, P. Wang, S. Wang, W. Wang, S. Wu, B. Xu, J. Xu, A. Yang, H. Yang, J. Yang, S. Yang, Y. Yao, B. Yu, H. Yuan, Z. Yuan, J. Zhang, X. Zhang, Y. Zhang, Z. Zhang, C. Zhou, J. Zhou, X. Zhou, T. Zhu, Qwen technical report, 2023. *arXiv preprint arXiv:2309.16609*.
- [61] Qwen2 technical report (2024).
- [62] H. Touvron, L. Martin, K. Stone, P. Albert, A. Almahairi, Y. Babaei, N. Bashlykov, S. Batra, P. Bhargava, S. Bhosale, D. Bikel, L. Blecher, C. Canton-Ferrer, M. Chen, G. Cucurull, D. Esiobu, J. Fernandes, J. Fu, W. Fu, B. Fuller, C. Gao, V. Goswami, N. Goyal, A. Hartshorn, S. Hosseini, R. Hou, H. Inan, M. Kardas, V. Kerkez, M. Khabsa, I. Kloumann, A. Korenev, P.S. Koura, M. Lachaux, T. Lavril, J. Lee, D. Liskovich, Y. Lu, Y. Mao, X. Martinet, T. Mihaylov, P. Mishra, I. Molybog, Y. Nie, A. Poulton, J. Reizenstein, R. Rungta, K. Saladi, A. Schelten, R. Silva, E.M. Smith, R. Subramanian, X.E. Tan, B. Tang, R. Taylor, A. Williams, J.X. Kuan, P. Xu, Z. Yan, I. Zarov, Y. Zhang, A. Fan, M. Kambadur, S. Narang, A. Rodriguez, R. Stojnic, S. Edunov, T. Scialom, Llama 2: open foundation and fine-tuned chat models, 2023. *arXiv preprint arXiv:2307.09288*.
- [63] S. Zhang, S. Roller, N. Goyal, M. Artetxe, M. Chen, S. Chen, C. Dewan, M. Diab, X. Li, X.V. Lin, T. Mihaylov, M. Ott, S. Shleifer, K. Shuster, D. Simig, P.S. Koura, A. Sridhar, T. Wang, L. Zettlemoyer, OPT: open pre-trained transformer language models, 2022. *arXiv:2205.01068*.
- [64] E. Almazrouei, H. Alobeidli, A. Alshamsi, A. Cappelli, R. Cojocaru, M. Debbah, E. Goffinet, D. Heslow, J. Launay, Q. Malartic, B. Noune, B. Pannier, G. Penedo, Falcon-40B: an open large language model with state-of-the-art performance (2023).