

Redefining Non-IID Data in Federated Learning for Computer Vision Tasks: Migrating From Labels to Embeddings for Task-Specific Data Distributions

Kasra Borazjani , *Student Member, IEEE*, Payam Abdisarabshali , *Student Member, IEEE*, Naji Khosravan ,
and Seyyedali Hosseinalipour , *Senior Member, IEEE*

Abstract—Federated learning (FL) has emerged as one of the prominent paradigms for distributed machine learning (ML). However, it is well-established that its performance can degrade significantly under non-IID (non-independent and identically distributed) data distributions across clients. To study this effect, the existing works predominantly emulate data heterogeneity by imposing label distribution skew across clients. In this article, we show that label distribution skew fails to fully capture the data heterogeneity in computer vision tasks beyond classification, *exposing an overlooked gap in the literature*. Motivated by this, by utilizing pretrained deep neural networks to extract task-specific data embeddings, we define *task-specific data heterogeneity* through the lens of each vision task and introduce a new level of data heterogeneity called *embedding-based data heterogeneity*. Our methodology involves clustering data points based on embeddings and distributing them among clients using the Dirichlet distribution. Through extensive experiments, we evaluate the performance of different FL methods under our revamped notion of data heterogeneity, introducing new benchmark performance measures to the literature. For instance, across seven representative computer vision tasks, our embedding-based heterogeneity formulation leads to up to around 60% increase in the observed loss under FedAvg, indicating that it more accurately exposes the performance degradation caused by data heterogeneity. We further unveil a series of open research directions that can be pursued.

Impact Statement—Federated learning (FL) is a cornerstone of privacy-preserving distributed machine learning (ML), yet its evaluation has been fundamentally shaped (and, as we show, misled) by the prevalent reliance on label-based data heterogeneity. Existing benchmarks overwhelmingly emulate non-IID (non-independent and identically distributed) data settings by

skewing label distributions, an approach suitable for classification but inadequate for vision tasks beyond classification (e.g., depth estimation, edge detection, and segmentation). Our work makes a conceptual and methodological leap by redefining data heterogeneity through task-specific embeddings. We introduce embedding-based data heterogeneity, extracted from the penultimate layer of task-trained deep neural networks, enabling the emulation of data heterogeneity from each vision task's perspective. Through extensive experiments across multiple FL algorithms and various representative computer vision tasks, we show that our embedding-based data heterogeneity produces significantly different performance trends, which reveals performance degradations invisible to label-based data heterogeneity models. As a result, this work establishes one of the first benchmarks for task-specific data heterogeneity, offers a new lens for studying task similarity, and opens fertile research directions in single- and multitask FL.

Index Terms—Computer vision, data heterogeneity, federated learning (FL), label distribution skew, non-IID data.

I. INTRODUCTION

FEDERATED learning (FL) has emerged as a pioneering paradigm for privacy-preserving distributed learning and has been applied across diverse domains such as vehicular networks [1], medical imaging [2], and sensing networks [3]. Compared to centralized training, a defining challenge in FL is that model performance is highly sensitive to data heterogeneity across participating devices/clients, often resulting in biased local model updates and degraded global convergence [4]. Consequently, understanding and accurately characterizing the impact of data heterogeneity has become a central research theme in FL [5]. In this article, we shed light on an underexplored dimension of data heterogeneity in FL by moving beyond the conventional *class-based formulations* (which rely on label distributions in classification tasks) to computer vision tasks such as “reshading,” where labels are inherently absent. This shift raises two fundamental questions: How does data heterogeneity impact the performance of FL in such vision tasks? And are current methodologies adequately capturing this impact? Our short answer is *No*. Specifically, we reveal a notable oversight in the literature—namely, that the performance of existing FL strategies in computer vision has been overestimated in the presence of data heterogeneity among the participating clients.

Received 2 August 2025; revised 30 November 2025, 28 December 2025, and 25 February 2026; accepted 19 March 2026. Date of publication 26 March 2026; date of current version 31 August 2026. This work was supported in part by the U.S. National Science Foundation (NSF) under Grant SaTC-2513164 and Grant ECCS-2512911. This article was recommended for publication by Associate Editor Guilherme Nelson DeSouza upon evaluation of the reviewers' comments. (*Corresponding author: Seyyedali Hosseinalipour.*)

Kasra Borazjani, Payam Abdisarabshali, and Seyyedali Hosseinalipour are with the Department of Electrical Engineering, University at Buffalo—SUNY, Buffalo, NY 14260 USA (e-mail: kasrabor@buffalo.edu; payamabd@buffalo.edu; alipour@buffalo.edu).

Naji Khosravan is with Adobe Firefly, Seattle, WA 98103 USA (e-mail: naji.khosravan@gmail.com).

This article has supplementary downloadable material available at <https://doi.org/10.1109/TAI.2026.3677838>, provided by the authors.

Digital Object Identifier 10.1109/TAI.2026.3677838

A. Overview and Literature Review

FL has emerged as a promising paradigm for distributed machine learning (ML), which enables multiple clients to collaboratively train ML models without sharing raw data, thus preserving user privacy [4], [6], [7]. This approach offers an alternative to traditional centralized ML techniques, where data must be transferred to a central server, often resulting in significant communication overhead and privacy concerns. The training process in FL takes place over several global rounds. In each round, clients first train ML models locally on their datasets, typically using several iterations of stochastic gradient descent (SGD). They then send the updated local models to a central server, which aggregates the models (usually through averaging) to form a new global model. This global model is subsequently broadcast back to the clients to start the next round of local training. Implementing FL in practice faces unique challenges, largely due to the *inherent non-IID (non-independent and identically distributed) or heterogeneous nature of data across clients* [6], [7], [8], [9], [10], [11], [12]. This data heterogeneity has a significant impact on FL performance. This challenge has motivated a range of FL methods designed to mitigate the adverse effects of data heterogeneity [7], [13], [14], [15].

Effective evaluation of the robustness of FL algorithms thus necessitates examining how data heterogeneity influences FL performance. To this end, state-of-the-art methods in FL that concern modeling the data heterogeneity have mainly emulated data heterogeneity across clients through imposing *label distribution skew*,¹ largely applied to computer vision *classification tasks* [6], [7], [9], [11], [16], [17], [18], [19], [20]. In particular, when FL first emerged as a decentralized ML paradigm, researchers emulated data heterogeneity through distributing data points based on their labels across clients to impose label distribution skew. This approach became a widely used method to simulate data heterogeneity in subsequent FL research [6], [11], [17], [19]. For example, in a handwritten digits classification task (MNIST dataset [21]), one client might have a local dataset mostly consisting of the digits “0” and “1,” while another client might only have images of the digits “7” and “8” [6]. As the number of labels per clients decreases, each client’s dataset becomes less representative of the global dataset, emulating a more non-IID data distribution. While this fundamental approach introduces some degree of data heterogeneity, there was a pressing need for mathematical and unified approaches to introduce and measure the level of data heterogeneity in FL. A major step in this direction was taken by utilizing the *Dirichlet distribution* to impose label distribution skew among clients in FL [7], [9], [11], [16], [18], [19], [20], enabling gradual heterogeneity in label/class distribution across the clients by fine-tuning Dirichlet concentration parameter.

Nevertheless, the above-mentioned class-based (equivalently, label-based) data heterogeneity may not be a perfect match

for other types of tasks, specifically computer vision tasks beyond classification, such as surface normal estimation, depth estimation, edge detection, and object detection [19], [22] (see Fig. 2 for a list of these tasks). Intuitively, the ineffectiveness of class-based data heterogeneity for generic computer vision tasks stems from the fact that non-classification vision tasks do not necessarily *perceive* data distribution in the same way as classification tasks. In fact, *their perception of data heterogeneity can be label-independent and majorly different from each other: an aspect largely overlooked in the FL literature*. For instance, the “2-D edges” task may be more sensitive to the configuration of edges within an image rather than its label. In this context, two data points with the same label (e.g., both with “Living Room” label) may exhibit highly different edge structures, rendering them heterogeneous from the perspective of the “2-D edges” task. Conversely, two data points with different labels (e.g., “Living Room” and “Kitchen”) might have similar edge structures, making them homogeneous for this task. This highlights the need for methodologies in FL that emulate data heterogeneity based on the unique perspectives of vision tasks to data.

Building on the above intuition, in this article, we propose the following fundamental hypothesis.

As we demonstrate in this article, the performance of FL for generic computer vision tasks has been overestimated in prior studies [23], [24], [25], [26]. These studies largely depend on class-based heterogeneity, which fails to accurately reflect the data distribution in non-classification computer vision tasks. Consequently, the actual performance and robustness of FL, and its variants, in handling data heterogeneity for generic computer vision tasks remain largely unexplored, highlighting a significant gap in the current literature.

To bridge this gap, we redefine the notion of data heterogeneity in FL for computer vision tasks as follows.

Instead of inducing heterogeneity solely from a raw data perspective, we propose a migration towards assessing data heterogeneity through the unique perspective of each task to the data after applying task-specific transformations on data.

To achieve the above objective, we introduce a new layer of data heterogeneity, which we call *embedding-based data heterogeneity*. To model *embedding-based data heterogeneity* (i.e., embedding distribution skew), we leverage a pretrained deep neural network (DNN) for each task and feed forward data points through this network (see Step 1 in Fig. 4). We then exploit the fact that the initial layers of the DNN of each task extract low-level features, and as the data progresses through deeper layers, the network extracts higher-level representations that are increasingly specific to the vision task of interest. Particularly, as we reach the layer just prior to the last layer (i.e., penultimate layer), *embeddings* that encapsulate a comprehensive view of the data from the perspective of the task are obtained (see Step 2 in Fig. 4), integrating both low-level and high-level features. We then leverage K-means to cluster data points based on these embeddings (see Step 3 in Fig. 4).

¹The label distribution skew approach involves partitioning data in such a way that clients receive data with imbalanced label distributions [4], [16].

Subsequently, we utilize a Dirichlet distribution to assign data points of each cluster to different clients to emulate task-specific data heterogeneity (see Step 4 in Fig. 4).

We note that, similar to the widely used class-based Dirichlet distribution approach, our embedding-based method is intended as a *data heterogeneity emulation mechanism* for the simulation phase of FL. In this setting, the method operates in a centralized manner on publicly available datasets prior to the start of FL training simulation to accommodate an offline generation of data partitions for clients in simulations. This preprocessing step offers a task-aware alternative to traditional class-based data splits and enables the study of how FL methods behave under the types of emulated data heterogeneity that can arise in non-classification computer vision tasks.

B. Summary of Contributions

Our main contributions can be summarized as follows.

- 1) We show through extensive experiments that prior FL studies relying on class/label distribution skew fail to impactfully model data heterogeneity for FL simulation in generic (non-classification) computer vision tasks.
- 2) We demonstrate that the notion of data heterogeneity is inherently task-dependent, motivating a shift away from the prevailing “one-size-fits-all” label-based formulation toward *task-specific data heterogeneity* models in FL.
- 3) We introduce *embedding-based data heterogeneity*, a new formulation that induces data heterogeneity from the perspective of each vision task, and show that FL performance can degrade under this revamped data heterogeneity model.
- 4) We outline a set of new research directions for handling task-specific data heterogeneity in single-task and multi-task FL settings involving computer vision tasks.

II. PRELIMINARIES AND MOTIVATIONS

In this section, we first model an FL system in Section II-A, describing the operations performed at both the client and server levels. We then discuss the current approaches for modeling data heterogeneity in FL in Section II-B, where we propose an experiment to evaluate these methods for computer vision tasks and highlight their limitations.

A. FL System Modeling

We consider executing FL over a set of clients denoted by $\mathcal{N}$. Each client $n \in \mathcal{N}$ is assumed to possess a local dataset $\mathcal{D}_n$ with size $D_n = |\mathcal{D}_n|$. We also let $\mathcal{D} = \bigcup_{n \in \mathcal{N}} \mathcal{D}_n$ with size $D = |\mathcal{D}| = \sum_{n \in \mathcal{N}} D_n$ denote the global dataset, encompassing the datasets distributed across the clients. For an ML task of interest, we let $\ell(\omega; d)$ denote the loss of the task under arbitrary ML model parameter ω and data point d . Subsequently, the local loss function of each client n is given by $\mathcal{L}_n(\omega) = 1/D_n \sum_{d \in \mathcal{D}_n} \ell(\omega; d)$. The goal of FL is to find the optimal model parameter ω^* to minimize the global loss

function $\mathcal{L}(\omega) = 1/D \sum_{d \in \mathcal{D}} \ell(\omega; d)$ for the ML task of interest, which can be mathematically represented as follows:

$$\omega^* = \arg \min_{\omega} \mathcal{L}(\omega) \equiv \arg \min_{\omega} \sum_{n \in \mathcal{N}} D_n \mathcal{L}_n(\omega) / D. \quad (1)$$

To achieve ω^* , FL follows an iterative approach consisting of a series of *local model training rounds* at the clients and *global model aggregation rounds* at the server. During each FL global aggregation round $t \in \{0, \dots, T-1\}$, the central server first broadcasts the global model, $\omega^{(t)}$, to all clients (with $\omega^{(0)}$ typically initialized randomly). Each client n will then initialize its local model as $\omega_n^{(t),0} = \omega^{(t)}$ and performs K rounds of local ML model training via mini-batch stochastic gradient descent (SGD), each indexed by k , where $1 \leq k \leq K$. In particular, the local model of client n after the k th SGD iteration during the global aggregation round t is given by

$$\omega_n^{(t),k} = \omega_n^{(t),k-1} - \eta_n^{(t),k} \tilde{\nabla} \mathcal{L}_n(\omega_n^{(t),k-1}) \quad (2)$$

where $\eta_n^{(t),k}$ is the instantaneous learning-rate. Also, $\tilde{\nabla}$ denotes the stochastic gradient approximation defined as $\tilde{\nabla} \mathcal{L}_n(\omega_n^{(t),k-1}) = 1/|\mathcal{B}_n^{(t),k-1}| \sum_{d \in \mathcal{B}_n^{(t),k-1}} \nabla \ell(\omega_n^{(t),k-1}; d)$, where $\mathcal{B}_n^{(t),k-1} \subseteq \mathcal{D}_n$ represents the set of randomly sampled data points contained in the mini-batch.

After the conclusion of local model training, clients send their latest trained local ML models (i.e., $\omega_n^{(t),K}$ for each client n) to the server, where the received models are aggregated to create a new global model as follows:

$$\omega^{(t+1)} = \sum_{n \in \mathcal{N}} D_n \omega_n^{(t),K} / D. \quad (3)$$

This new global model is then broadcast back to the clients to start the next local model training round and this iterative procedure is repeated until a desired convergence criterion is achieved. Note that the procedure formulated above describes the vanilla FedAvg method [6]. Other variations of FL, which are studied later in Section IV, involve modifications to the above-described local training and global aggregation processes.

B. Analyzing Current Approaches on Modeling Heterogeneity

1) *Importance and Impact of Data Heterogeneity in FL:* In FL, data is collected and stored locally by the clients, never being transferred over the network. Consequently, the server has often no control over the composition or distribution of the clients' local datasets. This leads to *data heterogeneity* (commonly referred to as non-IID data²) across clients, a fundamental challenge in FL. Specifically, when the datasets $\mathcal{D}_n$ and $\mathcal{D}_{n'}$ differ significantly across clients ($n \neq n'$), variations in local data distributions introduce *local model bias* during the local SGD iterations at the clients. This bias can degrade the overall performance of the FL global model. In particular, non-IID datasets result in local models that capture client-specific characteristics, which creates inconsistencies during model aggregation and exacerbates issues such as divergence and slow

²“non-IID” stands for “non-independent and identically distributed.”

convergence. *Modeling* and *addressing* data heterogeneity are therefore critical for advancing FL. In this work, we aim to introduce a novel *modeling* approach for non-IID data in FL, leaving *addressing* it to future research. Specifically, we focus on highlighting the inefficiencies of existing methods for modeling non-IID data in FL, particularly for computer vision tasks. Our proposed method provides a generic framework for data heterogeneity that not only extends to computer vision tasks but also encompasses existing methods as special cases.

2) *State-of-the-Art (SoTA) Methods in Modeling Data Heterogeneity*: SoTA methods in FL that concern modeling the data heterogeneity [11], [17], [18], [19] have heavily relied on classification tasks. In this setting, the global dataset often has a limited set of class labels, which streamlines the emulation of non-IID data across the clients through imposing imbalanced class label distribution across their local datasets. More trivial approaches [16], [27] include ensembling datasets to create a heterogeneous data gathering setting [27] and partially distributing labels between clients to emulate non-IID data (i.e., each client's local dataset contains data with a predetermined ratio from each class label) [16]. However, the most common approach is a theoretically-oriented method, which utilizes the *Dirichlet distribution* [9], [11], [16], [18], [19], [20] to emulate non-IID data. This SoTA method considers the ratio of each available class label in each client's local dataset as a random variable drawn from Dirichlet distribution which can be tuned with specific parameters. In the following, we first explain this SoTA method, which induces class-based data heterogeneity across the clients, and then study its effectiveness in more generic non-classification computer vision tasks.

3) *Creating Non-IID Data via Dirichlet Distribution*: To model the class-based data heterogeneity using the Dirichlet distribution in FL, a set of classification labels $\mathcal{C}$ with each member $c \in \mathcal{C}$ representing a class label in the global dataset is considered (e.g., each digit in MNIST [21] dataset). Then, parameter $\mathbf{p} = \{p_c\}_{c \in \mathcal{C}}$ is used for the set of ratios of each class c in the global dataset $\mathcal{D}$ which is being dispersed across clients. Afterwards, a label distribution vector is created for each client n as: $\mathbf{q}_n = \text{Dir}(\alpha\mathbf{p}) = \{q_{n,c}\}_{c \in \mathcal{C}}$ where $q_{n,c}$ denotes the ratio of client n 's dataset containing data from class label c . Here, $\alpha \geq 0$ is the *concentration parameter* which determines how closely each $\mathbf{q}_n$ will resemble $\mathbf{p}$, with higher α leading to a more similar $\mathbf{q}_n$ to $\mathbf{p}$. We can therefore interpret decreasing in the concentration parameter α as increasing the data heterogeneity across the clients. This phenomenon is visualized in Fig. 1 for an FL system with 25 clients containing 16 labels (similar results apply to other datasets such as MNIST [21], Fashion MNIST [28] and CIFAR-10 [29]). In the figure, data points with the same class label are shown in similar color tones. As illustrated, reducing α from 10^6 to 0.1 corresponds to increasing data heterogeneity, with $\alpha = 0.1$ representing an extreme case of non-IID data where most clients' datasets contain fewer than two distinct labels. As a result, it can be construed that smaller values of α represent emulation of scenarios where local datasets of clients induce severe local model bias during the local training rounds and negatively affect the convergence of FL.

Due to its reliance on class-label distributions (which are readily available for classification tasks) and its tunable concentration parameter α , the Dirichlet distribution has become the standard and widely adopted method for emulating non-IID data in FL [16], [20]. Specifically, its mathematical structure allows practitioners to precisely control data imbalance across clients, which explains its prominence as the de facto data heterogeneity model in the FL literature. Nevertheless, we next show that this method loses its effectiveness when FL is considered for more complex computer vision tasks.

4) *FL for non-classification Vision Tasks Under Class-Based Data Heterogeneity*: We next take one of the well-known computer vision datasets, called taskonomy [22], which has been previously used in both centralized [30] and federated [23] multitask learning. We utilize the above-described Dirichlet distribution to create a heterogeneous allocation of data across 25 clients' local datasets in FL using the available class labels, referred to as "class_scene" in this dataset.³ In our experiments, we select a representative set of vision tasks in this dataset, which are visualized in Fig. 2 and include "2-D edges," "reshading," "surface normals," "semantic segmentation," "3-D keypoints," "Euclidean depth estimation," and "scene classification" (unless otherwise stated, the loss functions used for the tasks are as follows: ℓ_1 loss for Euclidean depth estimation, 2-D edges, surface normals, and 3-D Keypoints; mean squared error for reshading; and cross-entropy for scene classification and semantic segmentation).

These vision tasks, except scene classification, can be construed as *non-classification tasks* since they are not associated with a class label. Following the non-IID data induced by Dirichlet distribution across the clients, we study the performance of the FL global model under different levels of heterogeneity (i.e., various α values) for the aforementioned vision tasks. The results of this experiment are shown in Fig. 3 (the left plot in each box, titled "class-based").

Focusing on the *only* classification task, i.e., scene classification (bottom box), we can see that decreasing α has a notable impact on the performance of the task. However, when focusing on all the other non-classification vision tasks, *we can see no significant degeneration in the performance of the global model as the value of α decreases*. This key observation signals the ineffectiveness of the class-based (equivalently, label-based) data heterogeneity emulation methods for generic computer vision tasks and calls for revamping the approaches used to induce data heterogeneity for such tasks in the FL setting.

Intuitively, class-based data heterogeneity emulation methods often fall short for generic computer vision tasks because non-classification vision tasks *perceive* data distribution differently from classification tasks. In fact, their perception of data heterogeneity can be *label-independent* and vary significantly across tasks, an aspect that has largely been overlooked in FL literature. This happens because labels provide only coarse

³The classes are provided in the dataset as probabilities for each data point belonging to each available class, from which we choose the class with the highest probability as the determined class for each data point. More details on the data processing are provided in Section IV-B.




Fig. 1. Dirichlet-generated data distribution for $N = 25$ clients and 16 classes with homogeneous prior probabilities (i.e., $p_i = 1/|\mathcal{C}|, \forall 1 \leq i \leq |\mathcal{C}|$). In the case of $\alpha = 10^6$, the data is homogeneously distributed across clients, whereas in the case of $\alpha = 0.1$, the majority of each client's data comprises two labels at most. The decrease in the homogeneity of the clients' data is also observed by comparing the data distributions generated by $\alpha = 10^3$ and $\alpha = 10$.




Fig. 2. Visualization of the tasks from the taskonomy dataset that have been used in our experiments. Unless otherwise stated, the loss functions used for the tasks are as follows: ℓ_1 loss for euclidean depth estimation, 2-D edges, surface normals, and 3-D keypoints; mean squared error for reshading; and cross-entropy for scene classification and semantic segmentation.

semantic groupings that do not necessarily reflect the underlying feature distributions relevant to each vision task. For example, consider the Euclidean depth estimation task, which focuses on estimating the depth of objects in an image. In this context, two images with the same label (e.g., both labeled “living room”) may exhibit highly different depth profiles due to variations in camera viewpoint or lighting conditions, making them heterogeneous from the perspective of this task. On the other hand, two images with different labels (e.g., “hallway” and “living room”) may share similar depth structures, making them homogeneous for this task despite their differing labels. This underscores the need for FL methodologies that account for data heterogeneity from the unique perspectives of each vision task, rather than relying solely on label-based notions of data heterogeneity.

Motivated by this observation, we propose an approach to emulate data heterogeneity in FL that accommodates both classification and non-classification computer vision tasks.

III. EMULATION OF DATA HETEROGENEITY FROM THE TASK PERSPECTIVE

To address the need for a task-specific modeling of data heterogeneity, we *shift our perspective* from the raw input domain to the *task output domain* in order to get more task-related information. In particular, instead of using the raw data features, such as its label, we aim to use the processed data features, such as *embeddings*, to emulate data heterogeneity in FL. This stems from the unavailability of representative characteristics, such as class labels, in the case of non-classification tasks, motivating us to identify a utility analogous to class labels that can effectively associate or disassociate data points based on the unique perspective of a given task.

In particular, we note that in the case of (deep) neural networks, training the network on different tasks leads to a re-configuration of its model parameters, where the parameters optimized for each task are often unique to that task. Consequently, feeding a raw data point into a neural network trained for a specific task and observing the outputs of different layers can provide unique representations of how the task perceives the data point. Specifically, the outputs of the initial layers, especially in case of convolutional neural networks (CNNs), often correspond to low-level features or simple patterns, while the outputs of the later layers correspond to high-level, task-specific features or other abstract representations. Following this notion, we aimed to extract the task's perspective of each data point by passing each data point through a network trained on the task and extracting the output of the penultimate layer as shown in Fig. 4. We will refer to this extracted output as *embedding* which we use to describe each data point from a task's perspective for the remainder of this discussion.

A. Task-Specific Data/Embedding Clustering

In a nutshell, in our approach (see Fig. 4), data heterogeneity is induced from the unique perspective of the task through an offline preprocessing stage, serving as a centralized client data generation procedure analogous to the class-based Dirichlet partitioning commonly adopted in prior work. First, data points are fed into a pretrained neural network trained specifically for the task of interest (any off-the-shelf pretrained model such as the ResNet [31] or CLIP [32] trained on the target task can be utilized). Next, embeddings for each data point are extracted from the penultimate layer of the model. These embeddings, which encapsulate task-specific features, are then clustered using algorithms such as K-means [33], [34] to reveal the similarity or dissimilarity of data points from the task's perspective (alternative choices to K-means have been studied



Fig. 3. *Class-based versus embedding-based distribution:* Comparison of how performing Dirichlet distribution over the data points' labels (equivalently, the scene class feature) and the extracted embeddings affect the performance of FL. Solid lines show the mean loss across clients; the shaded region denotes client-level dispersion (standard deviation of the loss across clients) at each global round. As it is shown, distributing the data points based on the class/label of the scene they correspond to (the left plot in each box) has not resulted in a significant change in the performance across various values of α . However, distributing the data points based on the clusters formed by the extracted embeddings (the right plot in each box) with the Dirichlet distributions of various α parameters has created a performance gap across the possible scenarios with $\alpha = 0.1$ (i.e., the most heterogeneous case) resulting in the worst performance and $\alpha = 1000$ (i.e., the least heterogeneous case) resulting in the best performance. The results further show that class-based data heterogeneity overestimates FL performance: it yields seemingly small loss values because it does not perturb the task-relevant feature space. In contrast, embedding-based heterogeneity increases the loss, revealing the sensitivity of FL methods to task-relevant feature variations that are not captured by class-based simulation splits in computer vision tasks. Also, it can be seen that only the class-based experiments on the scene classification task (i.e., the bottom-most box) exhibit a notable performance change with varying α , mirroring the trend seen in the embedding-based experiments. This indicates that the most salient feature for the scene classification task (i.e., the semantic class label) is naturally captured in the embeddings as well, leading to consistent changes in loss across both class-based and embedding-based settings.

in Appendix J). Finally, a Dirichlet distribution is applied to the clustered data to emulate the local datasets of clients, treating data points within the same cluster as belonging to the same *group* (analogous to “labels” in classification tasks).

Our method enables a task-driven and systematic approach to modeling data heterogeneity in FL, which works both for non-classification as well as classification tasks. In particular, in case of having a classification task, data points with the same label are expected to end up in the same cluster/group after

conducting the above procedure, which makes class-based data heterogeneity, explained in Section II-B, a special case of our method. We demonstrate this via visualizing our method on the scene classification task of taskonomy dataset, which is a classification task, in Fig. 5. As can be seen, the clustering of the data points with respect to their embeddings has resulted in the data points being clustered into separable clusters which possess only data points from one label. This mapping between the clustering of data points based on their embeddings and




Fig. 4. In our approach, data heterogeneity is induced from the unique task perspective. At first (i.e., in step 1), data points are fed to a pretrained neural network trained on the task (this can be any off-the-shelf model as long as the model is trained on the task). Afterwards (i.e., in step 2), the embeddings of data points are extracted from the penultimate layer of the model. Then (i.e., in step 3), the embeddings are clustered using clustering algorithms such as K-means, which unveils the similarity/dissimilarity of the data points from the task perspective. Finally (i.e., in step 4), Dirichlet distribution is applied on the clustered data to emulate the local datasets of clients, treating the data points in the same cluster as having the same *group* (analogous to *labels* in classification tasks).




Fig. 5. Figure illustrates embedding clusters extracted from a *scene classification* task model, where each color represents a specific class. It is evident that the embeddings create a separable space corresponding to the scene classification labels. This demonstrates that distributing data points across FL clients using the clusters/groups depicted in the figure is equivalent to distributing data based on the original scene class labels. The observed alignment between the embedding-based clustering and the class labels underscores the effectiveness of using embeddings as a method to induce task-specific data heterogeneity in FL.

their class labels for classification tasks further affirms that the clusters made from the embeddings can model the data points similarity/dissimilarity accurately from the task perspective and the clusters/groups obtained based on the embeddings can be hypothetically considered as having the same labels when applying the Dirichlet distribution.

Note that, for non-classification tasks, the clusters formed by the extracted embeddings may vary from those formed based on labels. In particular, we will later study the *overlap* between the clusters formed based on class labels (i.e., for the classification task) and those formed based on data embeddings of non-classification tasks (see Table I in Section IV).

B. Emulating Heterogeneous Local Datasets for Clients

Formally, once the clustering of data points based on their embeddings are obtained, we emulate the local datasets $\mathcal{D}_n$ for each client n , where the set of ratios for classes $\mathbf{p}$ detailed in Section II-B3 is now replaced by the distribution of the embeddings in each cluster/group, which we denote by $\hat{\mathbf{p}}$. Also,

we let the set $\mathcal{G}$ contain the distinct cluster/group indices (e.g., all 16 groups in Fig. 5). In our approach, each client n will be assigned data from each cluster/group based on its own cluster/group distribution vector $\hat{\mathbf{q}}_n = \text{Dir}(\alpha\hat{\mathbf{p}}) = \{\hat{q}_{n,g}\}_{g \in \mathcal{G}}$ where $\hat{q}_{n,g}$ denotes the ratio of client n 's dataset containing data from group g (e.g., one of the 16 groups in Fig. 5). These local datasets are then used during the local training of FL.

To demonstrate this revamped notion of data heterogeneity, we revisit the experiments in Section II-B to evaluate the effect of the induced data heterogeneity, replacing the class-based method with our embedding-based approach, which groups data points based on their embeddings. The results are presented in Fig. 3 (the right plot in each box, titled “embedding-based”) next to the class-based results for a better comparison. Unlike the findings from Section II-B, where data was distributed across clients based solely on class labels, we observe a pronounced impact of α (which controls the level of data heterogeneity across clients) on the performance of the trained FL global model. Specifically, as data heterogeneity increases (i.e., as α decreases), the global loss converges to a higher value, indicating a degradation in the model's performance.

These results highlight that task-specific heterogeneity is more effective in modeling a heterogeneous federated system for vision tasks compared to class-based heterogeneity, which has been predominantly used in prior FL studies [9], [10], [17], [18]. It is noteworthy that while we plot loss trajectories to visualize optimization behavior, another primary endpoint can be task-specific evaluation metrics computed at a fixed communication budget (20 global rounds). These standardized metrics are reported in Appendix D and summarized in Table IV in Appendix E for both data partitioning protocols as we will explain later. Furthermore, *the results reveal that the performance of FL for non-classification generic computer vision tasks has been overestimated in prior simulations [23], [24], [25], [26] as these works rely on class-based data heterogeneity*, which does not reflect the true data variations that vision tasks depend on (e.g., geometry, depth, texture, and illumination). In particular, when the data heterogeneity model fails to capture these task-relevant attributes, it leads to an overly optimistic assessment of FL performance under heterogeneous data conditions. Moreover, this experiment underscores that the actual performance and robustness of FL, and its variants, in handling data heterogeneity for generic computer vision tasks remain

TABLE I
HEATMAP REPRESENTATION OF TASK SIMILARITIES BASED ON THE MEAN ADJUSTED RAND INDEX (ARI) WITH PERMUTATION-BASED p -VALUES

Task 1 \ Task 2	Euclidean Depth Estimation	2-D Edges	3-D Keypoints	Surface Normals	Reshading	Semantic Segmentation	Scene Classification
Euclidean depth estimation	1.000	0.079 ($p=0.0099$)	0.044 ($p=0.0099$)	0.115 ($p=0.0099$)	0.127 ($p=0.0099$)	0.042 ($p=0.0099$)	0.031 ($p=0.0099$)
2-D edges	0.079 ($p=0.0099$)	1.000	0.060 ($p=0.0099$)	0.096 ($p=0.0099$)	0.066 ($p=0.0099$)	0.091 ($p=0.0099$)	0.043 ($p=0.0099$)
3-D keypoints	0.044 ($p=0.0099$)	0.060 ($p=0.0099$)	1.000	0.055 ($p=0.0099$)	0.041 ($p=0.0099$)	0.031 ($p=0.0099$)	0.045 ($p=0.0099$)
Surface normals	0.115 ($p=0.0099$)	0.096 ($p=0.0099$)	0.055 ($p=0.0099$)	1.000	0.102 ($p=0.0099$)	0.063 ($p=0.0099$)	0.040 ($p=0.0099$)
Reshading	0.127 ($p=0.0099$)	0.066 ($p=0.0099$)	0.041 ($p=0.0099$)	0.102 ($p=0.0099$)	1.000	0.054 ($p=0.0099$)	0.044 ($p=0.0099$)
Semantic segmentation	0.042 ($p=0.0099$)	0.091 ($p=0.0099$)	0.031 ($p=0.0099$)	0.063 ($p=0.0099$)	0.054 ($p=0.0099$)	1.000	0.061 ($p=0.0099$)
Scene classification	0.031 ($p=0.0099$)	0.043 ($p=0.0099$)	0.045 ($p=0.0099$)	0.040 ($p=0.0099$)	0.044 ($p=0.0099$)	0.061 ($p=0.0099$)	1.000

largely unknown—a critical gap in the literature. We address this gap through the benchmarks and experiments presented in Section IV.

In addition to the above-described main takeaways of the results, our approach has broader implications, which we elaborate on through the following remarks.

Remark 1 (Implications on Designing Effective Countermeasures to Address Data Heterogeneity in FL for Vision Tasks): Our work demonstrates that the prior notion of data heterogeneity, primarily based on the ratio of data points belonging to specific classes, is largely applicable only to classification tasks. By contrast, our method introduces a task-specific approach to inducing data heterogeneity, enabling experiments in FL to be extended to non-classification vision tasks. This shift not only broadens the scope of FL research to include a wider variety of vision tasks but also lays the groundwork for designing vision-specific methods to effectively handle and address/combat data heterogeneity in FL.

Remark 2 (Implications on the Concept of Task Similarity): Following the footsteps of prior work on modeling task similarity in the centralized ML domain [35], [36], our approach opens a model-driven perspective to quantify the similarity between two tasks. Specifically, this involves clustering data points based on their embeddings obtained from the models pre-trained for each task and comparing the resulting clusters to assess similarity. Using this straightforward approach, we establish a baseline performance for task similarity analysis, with results discussed in Section IV-F2. While this serves as a preliminary exploration, more sophisticated methods for task similarity measurement in the FL setting remain an open area for future investigation. Furthermore, this analysis provides an avenue for developing new task grouping strategies in FL, akin to prior efforts in centralized ML [37]. For example, by employing the task similarity measures based on the embedding-based clustering of data points, tasks with highly similar clustering patterns can be grouped and trained collaboratively, potentially enhancing FL performance for related tasks. We leave further exploration of this notion as future work.

Remark 3 (Implications on Multitask FL and Its Extensions): A fundamental implication of our method lies in its application to multitask FL and its extensions. In conventional multitask FL [25], [26], [38], each client typically has a single local task, with tasks varying across clients. On the other hand, the extended variations of multitask FL [23], [24] consider scenarios where multiple tasks coexist on a single FL client. In the context of conventional multitask FL, our method enables researchers to explore how varying levels of data heterogeneity, from the perspectives of different tasks, affect the performance of multitask FL systems. This opens the door to designing countermeasures specifically aimed at addressing task-induced data heterogeneity in multitask FL. Similarly, for the extensions of multitask FL, our method allows researchers to investigate how inducing data heterogeneity from the perspective of one task impacts the performance of other tasks trained concurrently on the same client. Moreover, it facilitates the study of grouping and training local tasks together under different levels of data heterogeneity, potentially leading to new strategies for optimizing task interactions in FL. We leave exploration of these concepts to future work.

Remark 4 (Complexity of Data Partitioning in FL): Configuring data heterogeneity across FL devices is inherently a combinatorial problem with an enormous solution space. Even in a simplified setting with 50,000 data points and 25 devices, the number of possible data splits is given by $\mathcal{S}(50,000, 25)$, where $\mathcal{S}(\cdot, \cdot)$ denotes the Stirling number of the second kind, which exceeds the estimated number of atoms in the universe ($\approx 10^{80}$). To address this, in a nutshell, our method introduces an ML-based preprocessing phase, which entails fine-tuning a pretrained model to extract embeddings for all data points (via one epoch of feedforward), and then clustering the embedded data. This process circumvents the brute-force nature of the combinatorial problem while enabling controlled configuration of data heterogeneity. We show that such clustering-based partitioning can produce client data distributions that impact global model performance, revealing a more severe performance gap than naive data splits. Importantly, the overhead of this preprocessing is justifiable as it is conducted entirely offline, before

any FL training. Moreover, in high-stakes scenarios (e.g., military and medical domains), investing in such a preprocessing step enables a more faithful stress-testing of FL methods under data heterogeneous conditions, rather than relying on benchmarks that may not fully capture task-relevant data variability.

IV. EXPERIMENTS AND BENCHMARKS

In this section, we experiment with various FL methods to give a performance benchmark for each method trained on each vision task under various levels of embedding-based heterogeneity. Our purpose in doing so is three-fold.

- 1) *Universality of the Phenomenon Across FL Methods*: We demonstrate that the phenomenon observed in Section III-B is universally applied across various FL methods and does not depend on the specific FL method employed.
- 2) *Consistency of Performance Impact on Vision Tasks*: We show the performance drop is consistent across all non-classification vision tasks.
- 3) *Benchmark for Future Research*: We provide a benchmark of reference for future endeavors to tackle the task-specific data heterogeneity across non-classification vision tasks in FL.⁴

A. Data Embedding Extraction

In our experiments, we use a pretrained CLIP [32] model which we fine-tune on the taskonomy dataset [22] for the tasks shown in Fig. 2. Then, we extract the embeddings from the fine-tuned model's penultimate layer and cluster them using a K-means algorithm with $K = 16$ for better comparison with class-based clusters (note that as we will discuss in Section IV-B, our dataset has 16 class labels). We have also provided the code on how we have fine-tuned the CLIP model in the above open-source GitHub repository which enables easy adaptations to other model architectures and datasets.

B. Data Preprocessing Details

The data gathered in the taskonomy dataset, pertaining to the scene classification label, is provided in the form of a probability vector $\beta_d = \{\beta_{d,1}, \beta_{d,2}, \dots, \beta_{d,J}\}$ for each data point d . Each $\beta_{d,j}$, $1 \leq j \leq J$ represents the probability of the data point d belonging to the scene class j in a set of $J = 365$ available classes. To simplify the process of distributing the data, we first take the max across the elements of each β_d for each data point and assign it as the class label to that data point. However, not all of the J classes have enough data points to be distributed across the simulated clients in our experiments: some classes have small numbers of data points which may fall as low as 1 data point. Thus, distributing the data points belonging to these ill-represented classes across the clients and including them in the classification task may lead to only some clients (as low as one) having data from these classes. This phenomenon will keep

us from ever having a homogeneous data distribution across the clients which is one of the settings we have experimented with throughout the article. Another problem is that training the global model in a federated fashion when one class is very poor in data point count will lead to an imbalanced training, where the biased classifying ability towards the ill-represented class will automatically worsen the results, rendering them unreliable for comparison across different hyperparameter values such as the Dirichlet distribution's concentration parameter α . To overcome these obstacles, we take the top 16 most-populated classes into consideration in all of our experiments and distribute their data across the clients.

C. Dataset Partitioning and Training Protocol

In the case of the Taskonomy dataset, for the centralized pretraining of each task, we use the initial three scenes ("allensville," "beechwood," and "benevolence") summing up to an initial count of 14 246 data points available. To track the performance of the centralized model, which is used for embedding extraction, for hyperparameter tuning and early stopping purposes, we use 100 data points chosen uniformly at random from the dataset as validation data. When performing the federated training, we distribute the data across the clients based on the underlying data partitioning method (i.e., embedding- and class-based), where 10% of local data points are used for validation. The numbers and statistics reported in the manuscript are based on this validation set. For global model validation, we report the average loss of the global model across the clients on their validation data. For the centralized training, we use the early stopping method to halt training if no improvement in loss was observed after five epochs. For the federated training, we train all methods for 20 global aggregations, each comprising 100 SGD rounds with the mini-batch size of 16. In addition, we sample all 25 clients during each global aggregation round as our method is not focused on client sampling. Unless stated otherwise, we fix random seeds for the FL training. To quantify uncertainty, we vary the K-means clustering seed and report mean $\pm$ std across seeds in Section IV-F4. For a summary of simulation hyperparameters refer to Table IV in Appendix E.

D. FL Methods

To observe the effect of embedding-based heterogeneity across various FL methods, we have considered the following six variations of FL.

- 1) *FedAvg* [6]: One of the widely used methods in FL with the dynamics explained in Section II-A. The results presented in Fig. 3 were also gathered by employing this method.
- 2) *FedProx* [7]: FedProx extends FedAvg by incorporating a proximal term into the local objective function to address data heterogeneity across clients. This term penalizes each client's local model divergence from the global model, encouraging consistency and stability during aggregation.
- 3) *SCAFFOLD* [13]: To address heterogeneous data distributions, SCAFFOLD maintains control variates, both on

⁴All of our implementations and results are contained in the public online GitHub repository of this article available at <https://github.com/KasraBorazjani/task-perspective-het.git>.

the server and clients, to adjust the local training updates and align them more closely with the global optimization direction. Clients send their updated control variates along with model parameters to the server. The server averages these control variates and redistributes them to clients in subsequent rounds.

- 4) *FedRep* [15]: FedRep decouples the optimization of shared and personalized components. Clients optimize a shared representation (encoder) collaboratively while maintaining a personalized classifier layer for task-specific outputs. During interaction with the server, only the shared representation is aggregated, leaving the personalized components untouched. This decoupling aims to minimize the negative impact of data heterogeneity while allowing personalization.
- 5) *FedAmp* [14]: FedAmp employs a model aggregation strategy that incorporates weighted averaging based on client-specific similarity measures. Clients send their local model updates and similarity scores to the server. The server aggregates these updates, prioritizing contributions from clients with higher similarity to the global objective. The method introduces an adaptive aggregation mechanism that accounts for the variability in clients' data distributions.

E. Adaptation of the Methods to Computer Vision Tasks

Unlike classification benchmarks, the taskonomy tasks used in this study (e.g., reshading and surface normals) require adapting FL baselines to encoder–decoder architectures with pixel-wise prediction heads. To ensure a fair comparison across methods, we adopt a consistent encoder–decoder backbone for all FL algorithms and specify, for each method, which components are treated as shared (aggregated) and which components are personalized below.

1) *Shared Versus Personalized Architecture Components*: For all methods, we decompose the network into an encoder that learns general feature representations and a decoder head that maps representations to pixel-wise or continuous outputs corresponding to the task being learned. The encoder corresponds to the aggregated component in methods such as FedAvg, FedProx, SCAFFOLD, and FedAmp. The decoder head is handled according to each method's personalization design, as detailed below.

2) *FedAvg, FedProx, and SCAFFOLD*: For these three baselines, the entire encoder–decoder network is treated as fully shared across all clients, and all parameters are aggregated after each communication round. FedProx additionally incorporates the proximal term during local updates, and SCAFFOLD employs control variates to correct local gradients. No structural changes to the architecture are required.

3) *FedRep*: FedRep's original formulation uses a personalized classifier layer. For pixel-wise tasks, we extend this idea by treating the decoder output layer as the personalized component while keeping the encoder and the rest of the decoder shared (i.e., aggregated across clients). This adaptation maintains the spirit of FedRep (i.e., personalizing the task-specific

output layer), while allowing consistent models to be learned collaboratively. In this regard, we follow the representation–head decoupling paradigm of [15], which emphasizes that the decoder or task head often captures client-specific variations in the downstream tasks. Our adaptations to FedRep adhere to this principle by personalizing decoder output layers where appropriate, while aggregating the rest of the model parameters.

4) *FedAmp*: We make two modifications to adapt FedAmp to the pixel-wise prediction tasks in taskonomy. First, we include the full model parameters (encoder + decoder) in the attentive message passed between clients during the proximal update. Second, to compute client–client similarity, we measure the distance between their decoder parameters (reflecting task-specific behavior) and apply softmax normalization to obtain similarity weights for aggregation.

5) *Loss Functions*: The vision task's loss functions used across all methods are identical (represented in the y-axis of the figures). All extra loss terms correspond to the original FL method (e.g., for FedProx and FedAmp).

6) *Parameter Budget Alignment*: All FL methods use the same encoder–decoder backbone, ensuring identical parameter counts (methods that personalize the decoder (e.g., FedRep) add no extra shared parameters). Thus, performance differences reflect method behavior (not architectural advantage).

F. Further Results and Discussions

1) *Studying the Performance of Various Benchmarks Under Different Levels of Heterogeneity*: To provide a benchmark of FL performance, we compare the loss value of the benchmarks⁵ explained above under class-based and embedding-based data distribution (the same system of 25 clients as in Section II-B4 is considered). This benchmarking provides a baseline performance for each of these SoTA methods under various levels of data heterogeneity. Note that we have shown the results for FedAvg in Fig. 3, which were used to explain our methodology. Subsequently, in the following, we discuss the results for FedProx (see Fig. 6), SCAFFOLD (see Fig. 7), FedRep (see Fig. 8), and FedAmp (see Fig. 9).

Observing the results in Figs. 6–9, we can observe the same phenomenon as in Fig. 3 (see Section III-B). In particular, the results indicate that, for the non-classification tasks (i.e., all the tasks except “scene classification” presented in the bottom box of each figure), the change in the value of α in the Dirichlet distribution in all methods results in a performance gap when our embedding-based distribution of data is utilized (see the right plot in each box). However, at the same time, it can be observed that such a performance gap does not hold for class-based distribution of data (see the left plot in each box). Further, our results verify that embedding-based data heterogeneity and class-based data heterogeneity result in a similar FL performance for the classification task (see the bottom box in Figs. 6–9). The main two takeaways of the results are as follows: a) these observations further validate that our proposed method of

⁵Note that “loss” is used as a unified metric that enables universal comparisons across the tasks. For example, other metrics such as *accuracy*, which is mostly tailored to classification, do not benefit from such a universality.




Fig. 6. *Class-based versus embedding-based distribution*: Comparison of how performing Dirichlet distribution over the data points' labels (equivalently, the scene class feature) and the extracted embeddings affect the performance in FL when the FedProx technique is used. Solid lines show the mean loss across clients; the shaded region denotes client-level dispersion (standard deviation of the loss across clients) at each global round. The same phenomena observed in Fig. 3 can be observed in the FedProx results.

inducing heterogeneity from the tasks' perspective is capable of creating a heterogeneous data setting more effectively for non-classification vision tasks; and b) in addition to revealing the performance of all the benchmarks under embedding-based data heterogeneity, the results reveal that all these SoTA methods are subject to overestimation of FL performance under class-based data heterogeneity for vision tasks: when class-based data heterogeneity is induced, the loss is often lower than that of the embedding-based data heterogeneity scenario. In other words, embedding-based data heterogeneity can induce a more pronounced impact on the FL performance (i.e., a higher value of loss).

Further, a major takeaway of implementing these techniques were as follows: the tasks exhibit model bias at different numbers of local SGD iterations (K), which depends on the task complexity and the underlying data distributions. In other

words, some tasks may need to be trained over longer durations of local training (i.e., higher values of K) to exhibit the model bias and subsequently a gap in performance under various levels of data heterogeneity. Nevertheless, to have fair comparisons between class-based and embedding-based methods, we maintained the same number of local SGDs per round for both class-based and embedding-based experiments across all benchmark methods and all tasks. The details of these hyperparameters are available at our aforementioned GitHub repository. Also, for an at-a-glance comparison of all FL methods, refer to Appendix A.

2) *Studying the Task Similarity Through Comparing the Embedding Clusters*: Revisiting Remarks 2 and 3, it can be noted that the extracted embeddings of the data points and the way that they get clustered in our proposed approach may have further use cases. As a preliminary investigation in this direction, in Table I, we calculate the similarity of embedding




Fig. 7. *Class-based versus embedding-based distribution*: Comparison of how performing Dirichlet distribution over the data points' labels (equivalently, the scene class feature) and the extracted embeddings affect the performance in FL when the SCAFFOLD technique is used. Solid lines show the mean loss across clients; the shaded region denotes client-level dispersion (standard deviation of the loss across clients) at each global round. The same phenomena observed in Fig. 3 can be observed in the SCAFFOLD results.

clusters across pairs of tasks through *adjusted rand index (ARI)*, a standard clustering measure with well-understood statistical properties [39], and visualize the resulting cross-task similarities using a heatmap. In essence, ARI measures how consistently two different tasks cluster the same set of data points: it counts how often pairs of data samples are placed in the same cluster or in different clusters by both tasks, and then corrects this agreement for what would be expected purely by chance. As a result, ARI takes the value 1 when the two tasks clusterings match perfectly, 0 when the alignment is no better than random, and can even become negative when the agreement is worse than random.

Inspecting the results in Table I, the *Euclidean depth estimation* and the *reshading* tasks have the highest average ARI across their embedding clusters (0.127) among the pairs included in the table. The similarity between the output domains

of these two tasks and the general aim each one is accomplishing can help explain the higher score: both of these tasks are driven by the same geometric structure of a scene (e.g., how surfaces are shaped, oriented, and spatially arranged) and since their outputs rely on these shared cues, they partition images along similar geometric patterns, resulting in more closely aligned embedding clusters and a higher ARI. On the other hand, the two lowest similarity ratios (i.e., both 0.031) belong to the *3-D Keypoints* task when compared with *semantic segmentation* task and *Euclidean depth estimation* when compared with *scene classification* task. This can also be explained through the fact that both 3-D keypoints and Euclidean depth estimation are tasks that consider the very details of the image which is contrary to the higher-level inspection conducted in the semantic segmentation and scene classification tasks. The rest of the pairwise comparisons fall in between these two extremes.



Fig. 8. *Class-based versus embedding-based distribution*: Comparison of how performing Dirichlet distribution over the data points' labels (equivalently, the scene class feature) and the extracted embeddings affect the performance in FL when the FedRep technique is used. Solid lines show the mean loss across clients; the shaded region denotes client-level dispersion (standard deviation of the loss across clients) at each global round. The same phenomena observed in Fig. 3 can be observed in the FedRep results.

In particular, many of the numerical values are hard to interpret intuitively since human-level visual explanations may not suffice to describe the innate characteristics of these tasks. This makes the reach of the methodology used to obtain the heatmap in Table I more pronouncedly felt in scenarios that human-level intuitions are nonexistent. The extensions of the implications of these results to task grouping and multitask training is left as future work (see Remarks 2 and 3). Further, it can be noted that the ARI values reported in Table I are numerically small (reflecting that the tasks are largely distinct from one another), where the accompanying permutation-based p-values confirm that these similarities are statistically meaningful/reliable. In particular, the consistently low p-values ($p = 0.0099 \simeq 0.01$, measured across 100 permutations) show that even the modest ARI scores observed are significantly above what would arise by chance, indicating task-induced structure rather than noise.

3) *Studying the Effect of the Number of Clusters on the Silhouette Score*: To assess the structural quality of the learned embeddings, we employ the *Silhouette coefficient*, a standard cluster validity index [40]. In words, Silhouette coefficient measures how well each data point fits within its assigned cluster compared to its proximity to neighboring clusters, summarizing both compactness and separation. Its values range in $[-1, 1]$, where higher values indicate clearer and more meaningful cluster structure. We compute this coefficient for $k \in \{2, 4, 10, 16, 32\}$ across all tasks. As shown in Appendix B, the value of Silhouette score increases marginally from $K = 16$ onward. This diminishing return, combined with the increase in computational cost associated with larger k , motivates our choice of $n_{\text{clusters}} = 16$ for all analyses in this article.

4) *Studying the Effect of Randomness in Clustering on the Stability of the Observed Gaps in Performance*: We analyze




Fig. 9. *Class-based versus embedding-based distribution*: Comparison of how performing Dirichlet distribution over the data points' labels (equivalently, the scene class feature) and the extracted embeddings affect the performance in FL when the FedAmp technique is used. Solid lines show the mean loss across clients; the shaded region denotes client-level dispersion (standard deviation of the loss across clients) at each global round. The same phenomena observed in Fig. 3 can be observed in the FedAmp results.

the effect of different random seeds used in K-means clustering on the performance gaps observed under our embedding-based data heterogeneity formulation. To this end, we revisit our main results in Fig. 3 and repeat the experiments using three distinct K-means seeds. We then report the mean of the resulting loss values alongside their uncertainty intervals (standard deviation) for each α in Table II. The results show that the performance degradation at lower α values (i.e., higher loss values) consistently persists even after averaging across seeds, demonstrating that our embedding-based heterogeneity remains effective regardless of the seed initialization. Moreover, the standard deviations are significantly smaller than the corresponding mean losses, indicating that the observed performance gaps are stable and not driven by seed-dependent fluctuations: intuitively, this robustness arises because the embedding space already encodes strong geometric and structural cues, causing

K-means to produce similar cluster partitions across different initializations.

5) *Studying Task Performance on Another Dataset*: To evaluate the generality of our approach beyond taskonomy, we further compare embedding-based and class-based data heterogeneity on a different vision dataset: PASCAL [41], a widely used dataset spanning diverse object categories and scene structures. For this study, we apply FedAvg and report the results in Appendix C. Consistent with our main findings, the embedding-based data splits on PASCAL induce clear performance gaps across various data heterogeneity levels, whereas the class-based data splits fail to meaningfully perturb the task-relevant feature space and have only a negligible effect on the trained FL model's performance.

6) *Studying Task-Specific Metrics*: As mentioned earlier, our choice of loss functions (namely, ℓ_1 loss for Euclidean

TABLE II
AVERAGE AND STANDARD DEVIATION OF THE LOSS FOR EACH TASK IN TASKONOMY USING FEDAVG, ACROSS RANDOM SEEDS OF K-MEANS

α	Tasks						
	2-D Edges	Reshading	Surface Normals	Semantic Segmentation	3-D Keypoints	Euclidean Depth Estimation	Scene Classification
0.1	0.077 ± 0.002	0.057 ± 0.001	0.196 ± 0.005	0.052 ± 0.001	0.090 ± 0.002	0.075 ± 0.002	2.324 ± 0.065
10	0.074 ± 0.002	0.054 ± 0.001	0.190 ± 0.004	0.048 ± 0.001	0.085 ± 0.002	0.074 ± 0.002	1.779 ± 0.053
1000	0.072 ± 0.002	0.048 ± 0.001	0.186 ± 0.005	0.048 ± 0.001	0.084 ± 0.002	0.071 ± 0.002	1.357 ± 0.033

depth estimation, 2-D edges, surface normals, and 3-D keypoints; mean squared error for reshading; and cross-entropy for scene classification and semantic segmentation) was motivated by their generic applicability across a broad range of vision tasks. For concreteness, we additionally provide a study using more task-specific metrics, where we employ F-measure for 2-D edges, PSNR for reshading, mIoU for semantic segmentation, RMSE for Euclidean depth estimation, and accuracy for scene classification. The results of these experiments are presented in Appendix D. Consistent with our main findings, the embedding-based data heterogeneity again produces clear performance gaps across different heterogeneity levels, whereas class-based heterogeneity yields negligible performance variation.

7) *Studying other Embedding Extraction Methods:* We conduct an ablation study comparing generic and task-aligned embedding extraction strategies, including frozen CLIP, self-supervised SimCLR, and task-supervised CLIP, to examine their ability to induce task-perspective data heterogeneity. The study shows that only task-supervised embeddings consistently reveal clear performance separation between low- and high-heterogeneity regimes, while generic embeddings fail to do so. Detailed results and analysis are provided in Appendix F.

8) *Studying the Effect of Embedding Extraction Layer:* We conduct an ablation study comparing embeddings extracted from multiple decoder layers to assess sensitivity to the extraction point. The qualitative separation between low- and high-heterogeneity regimes persists across all layers, while the penultimate layer yields slightly larger performance gaps under severe heterogeneity. Full results are reported in Appendix G.

9) *Effect of Training Backbone on Embedding-Based Heterogeneity:* We further assess robustness to architectural coupling by training the federated global model with a ResNet-50 backbone while retaining client partitions generated from task-supervised CLIP embeddings. The embedding-based heterogeneity effect persists despite this backbone change, indicating that the observed performance degradation is not an artifact of encoder alignment. Full results are reported in Appendix H.

10) *Effect of the Number of Clusters k :* We further analyze the sensitivity of the proposed method to the number of clusters k used during embedding-based partitioning. We vary k over a meaningful range and observe that, for moderate to large values, the monotonic relationship between data heterogeneity and performance and the induced separation between low- and

high-heterogeneity regimes are preserved. Detailed results and analysis are provided in Appendix I.

11) *Qualitative Behavior of Client Data Partitions Under Varying K :* Although we use $k = 16$ as the default number of embedding clusters, we additionally evaluate the effect of varying $k \in \{2, 4, 10, 16, 32\}$. As detailed in Appendix K, increasing k refines the granularity of embedding partitions without introducing drastic cluster imbalance or effectively alter client sample compositions. Across all tested values of k , the qualitative heterogeneity structure induced by the Dirichlet distribution is preserved, which in turn explains the observed stability of the monotonic α -performance trends.

V. CONCLUSION AND FUTURE WORK

In this article, we demonstrated that inducing data heterogeneity in FL via a label/class-based approach alone fails to fully encapsulate the data heterogeneity encountered in computer vision tasks beyond classification. To address this limitation, we extracted task-specific data embeddings and proposed a new perspective on data heterogeneity in FL: *embedding-based data heterogeneity*. By clustering data points based on their embeddings and distributing them among clients using the Dirichlet distribution, we introduce a more nuanced formulation of data heterogeneity in FL, particularly for non-classification vision tasks. Our experiments validated the impact of this revised framework, providing new benchmark performance measures and offering various nuanced insights. Beyond these contributions, we highlighted several open research directions, paving the way for future advancements in FL under this revised notion of data heterogeneity.

Further, it is worth mentioning that although our embedding extraction is performed centrally for benchmarking, consistent with standard Dirichlet-based emulation, future work may incorporate secure aggregation protocols [42] to enable clients to compute embeddings locally and share only aggregated or privatized statistics to the server to obtain global knowledge about data heterogeneity across the clients. Finally, while our current study treats the vision tasks as a flat collection, future extensions of our benchmark could incorporate hierarchical structure and uncertainty-aware evaluation, inspired by [43]. For example, related tasks (e.g., depth, surface normals, and reshading) could be grouped into higher-level families, allowing data heterogeneity trends to be analyzed not only per task but also at the level of task families.

REFERENCES

- [1] S. A. Abdel Hakeem and H. Kim, "Advancing intrusion detection in V2X networks: A comprehensive survey on machine learning, federated learning, and edge AI for V2X security," *IEEE Trans. Intell. Transp. Syst.*, vol. 26, no. 8, pp. 11137–11205, Aug. 2025.
- [2] K. Borazjani, N. Khosravan, L. Ying, and S. Hosseinalipour, "Multi-modal federated learning for cancer staging over non-IID datasets with unbalanced modalities," *IEEE Trans. Med. Imag.*, vol. 44, no. 1, pp. 556–573, Jan. 2025.
- [3] A. Kapoor and D. Kumar, "Federated learning for urban sensing systems: A comprehensive survey on attacks, defences, incentive mechanisms, and applications," *IEEE Commun. Surveys Tuts.*, vol. 27, no. 2, pp. 1293–1325, Feb. 2025.
- [4] P. Kairouz et al., "Advances and open problems in federated learning," *Foundations Trends Mach. Learn.*, vol. 14, no. 1, pp. 1–210, 2021.

- [5] J. Pei, W. Liu, J. Li, L. Wang, and C. Liu, "A review of federated learning methods in heterogeneous scenarios," *IEEE Trans. Consum. Electron.*, vol. 70, no. 3, pp. 5983–5999, Mar. 2024.
- [6] B. McMahan, E. Moore, D. Ramage, S. Hampson, and B. A. y. Arcas, "Communication-efficient learning of deep networks from decentralized data," in *Proc. Int. Conf. Artif. Intell. Statist.*, vol. 54. PMLR, Apr. 2017, pp. 1273–1282.
- [7] T. Li, A. K. Sahu, M. Zaheer, M. Sanjabi, A. Talwalkar, and V. Smith, "Federated optimization in heterogeneous networks," *Proc. Mach. Learn. Syst. (MLSys)*, vol. 2, pp. 429–450, 2020.
- [8] S. Pieri, J. Restom, S. Horvath, and H. Cholakkal, "Handling data heterogeneity via architectural design for federated visual recognition," *Adv. Neural Inf. Process. Syst.*, vol. 36, pp. 4115–4136, 2023.
- [9] M. Mendieta, T. Yang, P. Wang, M. Lee, Z. Ding, and C. Chen, "Local learning matters: Rethinking data heterogeneity in federated learning," in *Proc. IEEE/CVF Conf. Comput. Vision Pattern Recognit.*, 2022, pp. 8397–8406.
- [10] M. Luo, F. Chen, D. Hu, Y. Zhang, J. Liang, and J. Feng, "No fear of heterogeneity: Classifier calibration for federated learning with non-IID data," in *Advances in Neural Information Processing System*, vol. 34. Curran Associates Inc., 2021, pp. 5972–5984.
- [11] E. Darzi, Y. Shen, Y. Ou, N. M. Sijtsma, and P. M. van Ooijen, "Tackling heterogeneity in medical federated learning via aligning vision transformers," *Artif. Intell. Med.*, vol. 155, Art. no. 102936, 2024.
- [12] M. Zhang, L. Qu, P. Singh, J. Kalpathy-Cramer, and D. L. Rubin, "SplitAVG: A heterogeneity-aware federated deep learning method for medical imaging," *IEEE J. Biomed. Health Inform.*, vol. 26, no. 9, pp. 4635–4644, Sep. 2022.
- [13] S. P. Karimireddy, S. Kale, M. Mohri, S. Reddi, S. Stich, and A. T. Suresh, "SCAFFOLD: Stochastic controlled averaging for federated learning," in *Proc. Int. Conf. Machine Learn.* PMLR, 2020, pp. 5132–5143.
- [14] Y. Huang et al., "Personalized cross-silo federated learning on non-iid data," in *Proc. AAAI Conf. Artif. Intell.*, vol. 35, no. 9, pp. 7865–7873, 2021.
- [15] L. Collins, H. Hassani, A. Mokhtari, and S. Shakkottai, "Exploiting shared representations for personalized federated learning," in *Proc. Int. Conf. Mach. Learn.*, vol. 139. PMLR, 18–24 Jul 2021, pp. 2089–2099.
- [16] Q. Li, Y. Diao, Q. Chen, and B. He, "Federated learning on non-iid data silos: An experimental study," in *IEEE Int. Conf. Data Eng. (ICDE)*. Piscataway, NJ, USA: IEEE Press, 2022, pp. 965–978.
- [17] L. Qu et al., "Rethinking architecture design for tackling data heterogeneity in federated learning," in *Proc. IEEE/CVF Conf. Comput. Vision Pattern Recognit.*, 2022, pp. 10061–10071.
- [18] T.-M. H. Hsu, H. Qi, and M. Brown, "Measuring the effects of non-identical data distribution for federated visual classification," 2019, *arXiv:1909.06335*.
- [19] C. He et al., "FedCV: A federated learning framework for diverse computer vision tasks," 2021, *arXiv:2111.11066*.
- [20] J. Wang, Q. Liu, H. Liang, G. Joshi, and H. V. Poor, "Tackling the objective inconsistency problem in heterogeneous federated optimization," *Adv. Neural Inf. Process. Syst.*, vol. 33, pp. 7611–7623, 2020.
- [21] Y. Lecun, L. Bottou, Y. Bengio, and P. Haffner, "Gradient-based learning applied to document recognition," *Proc. IEEE*, vol. 86, no. 11, pp. 2278–2324, 1998.
- [22] A. R. Zamir, A. Sax, W. Shen, L. J. Guibas, J. Malik, and S. Savarese, "Taskonomy: Disentangling task transfer learning," in *Proc. IEEE Conf. Comput. Vision Pattern Recognit.*, 2018, pp. 3712–3722.
- [23] W. Zhuang, Y. Wen, L. Lyu, and S. Zhang, "MAS: Towards resource-efficient federated multiple-task learning," in *Proc. IEEE/CVF Int. Conf. Comput. Vis.*, 2023, pp. 23414–23424.
- [24] Y. Lu, S. Huang, Y. Yang, S. Sirejiding, Y. Ding, and H. Lu, "FedHCA2: Towards hetero-client federated multi-task learning," in *Proc. IEEE/CVF Conf. Comput. Vision Pattern Recognit.*, 2024, pp. 5599–5609.
- [25] M. Mortaheb, C. Vahapoglu, and S. Ulukus, "Personalized federated multi-task learning over wireless fading channels," *Algorithms*, vol. 15, no. 11, 2022, Art. no. 421.
- [26] R. Cai et al., "Many-task federated learning: A new problem setting and a simple baseline," in *Proc. IEEE/CVF Conf. Comput. Vision Pattern Recognit.*, 2023, pp. 5037–5045.
- [27] H. Chen, Y. Zhang, D. Krompass, J. Gu, and V. Tresp, "FedDat: An approach for foundation model finetuning in multi-modal heterogeneous federated learning," in *Proc. AAAI Conf. Artif. Intell.*, vol. 38, no. 10, pp. 11285–11293, 2024.
- [28] H. Xiao, K. Rasul, and R. Vollgraf, "Fashion-MNIST," 2025. [Online]. Available: <https://github.com/zalando-research/fashion-mnist>
- [29] A. Krizhevsky and G. Hinton, "Learning multiple layers of features from tiny images," 2009, *arXiv:1826.8744*.
- [30] A. Sax, B. Emi, A. R. Zamir, L. Guibas, S. Savarese, and J. Malik, "Mid-level visual representations improve generalization and sample efficiency for learning visuomotor policies," 2018, *arXiv:1812.11971*.
- [31] K. He, X. Zhang, S. Ren, and J. Sun, "Deep residual learning for image recognition," in *Proc. IEEE Conf. Comput. Vision Pattern Recognit.*, 2016, pp. 770–778.
- [32] A. Radford et al., "Learning transferable visual models from natural language supervision," in *Proc. Int. Conf. Mach. Learn.* PMLR, 2021, pp. 8748–8763.
- [33] S. Lloyd, "Least squares quantization in PCM," *IEEE Trans. Inf. Theory*, vol. 28, no. 2, pp. 129–137, Feb. 1982.
- [34] J. B. McQueen, "Some methods of classification and analysis of multivariate observations," in *Proc. 5th Berkeley Symp. Math. Stat. Prob.*, 1967, pp. 281–297.
- [35] C. Shui, M. Abbasi, L.-É. Robitaille, B. Wang, and C. Gagné, "A principled approach for learning task similarity in multitask learning," 2019, *arXiv:1903.09109*.
- [36] C. P. Le, M. Soltani, J. Dong, and V. Tarokh, "Fisher task distance and its application in neural architecture search," *IEEE Access*, vol. 10, pp. 47235–47249, 2022.
- [37] T. Standley, A. Zamir, D. Chen, L. Guibas, J. Malik, and S. Savarese, "Which tasks should be learned together in multi-task learning?," in *Int. Conf. Mach. Learn.* PMLR, 2020, pp. 9120–9132.
- [38] V. Smith, C.-K. Chiang, M. Sanjabi, and A. S. Talwalkar, "Federated multi-task learning," *Adv. Neural Inf. Process. Syst.*, vol. 30, pp. 234–245, 2017.
- [39] L. Hubert and P. Arabie, "Comparing partitions," *J. Classification*, vol. 2, no. 1, pp. 193–218, 1985.
- [40] K. R. Shahapure and C. Nicholas, "Cluster quality analysis using silhouette score," in *Int. Conf. Data Sci. Adv. Anal. (DSAA)*. Piscataway, NJ, USA: IEEE Press, 2020, pp. 747–748.
- [41] R. Mottaghi et al., "The role of context for object detection and semantic segmentation in the wild," in *Proc. IEEE Conf. Comput. Vision Pattern Recognit.*, 2014, pp. 891–898.
- [42] K. Bonawitz et al., "Practical secure aggregation for privacy-preserving machine learning," in *Proc. ACM SIGSAC Conf. Comput. Commun. Secur.*, 2017, pp. 1175–1191.
- [43] O. Barkan, J. Benichimol, I. Caspi, E. Cohen, A. Hammer, and N. Koenigstein, "Forecasting CPI inflation components with hierarchical recurrent neural networks," *Int. J. Forecasting*, vol. 39, no. 3, pp. 1145–1162, 2023.
- [44] T. Chen, S. Kornblith, M. Norouzi, and G. Hinton, "A simple framework for contrastive learning of visual representations," in *Proc. Int. Conf. Mach. Learn.* PMLR, 2020, pp. 1597–1607.

Kasra Borazjani (Student Member, IEEE) is currently working toward the Ph.D. degree in electrical engineering with the University at Buffalo-SUNY, Buffalo, NY, USA.

Payam Abdisarabshali (Student Member, IEEE) is currently working toward the Ph.D. degree in electrical engineering with the University at Buffalo-SUNY, Buffalo, NY, USA.

Naji Khosravan received the Ph.D. degree from the University of Central Florida in 2019.
He is a Senior Manager at Adobe Firefly, Seattle, WA USA.

Seyyedali Hosseinalipour (Student Member, IEEE) received the Ph.D. degree from the NC State University in 2020.
He is an Assistant Professor of electrical engineering with the Department at the University at Buffalo-SUNY, Buffalo, NY, USA.